

Georgiy_Yudintsev_Data_Challenge_1

February 20, 2019

```
In [300]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

Read in the data frame.

```
In [301]: Empl_retten = pd.read_csv("employee_retention_data.csv")
Empl_retten.head(5)
```

```
Out[301]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	

	quit_date
0	2015-10-30
1	2014-04-04
2	NaN
3	2013-06-07
4	2014-08-22

The next couple commands are to create new variables for duration: duration of employees still working (up to '2015-12-13') and duration of employees who quit.

```
In [302]: Empl_retten['still_working'] = np.where(pd.isnull(Empl_retten['quit_date']), ['2015-12-13'], Empl_retten['quit_date'])
```

```
In [303]: Empl_retten['join_date'] = pd.to_datetime(Empl_retten['join_date'])
Empl_retten['quit_date'] = pd.to_datetime(Empl_retten['quit_date'])
Empl_retten['still_working'] = pd.to_datetime(Empl_retten['still_working'])
```

```
In [304]: Empl_retten.head(5)
```

```
Out[304]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	

2	927315.0	4	marketing	14	101000.0	2014-10-13
3	662910.0	7	customer_service	20	115000.0	2012-05-14
4	256971.0	2	data_science	23	276000.0	2011-10-17

	quit_date	still_working
0	2015-10-30	NaT
1	2014-04-04	NaT
2	NaT	2015-12-13
3	2013-06-07	NaT
4	2014-08-22	NaT

Compute the duration of employment for employees who quit. Add this variable as a new column to the existing dataframe.

```
In [305]: Empl_reten['empl_quit_duration'] = (Empl_reten['quit_date'] - Empl_reten['join_date']).dt.days
```

Compute the duration of employment for employees who are still working (up to '2015-12-13'). Add this variable as a new column to the existing dataframe.

```
In [306]: Empl_reten['still_working_duration'] = (Empl_reten['still_working'] - Empl_reten['join_date']).dt.days
```

```
In [307]: Empl_reten.head(5)
```

```
Out[307]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	


	quit_date	still_working	empl_quit_duration	still_working_duration
0	2015-10-30	NaT	585 days	NaT
1	2014-04-04	NaT	340 days	NaT
2	NaT	2015-12-13	NaT	426 days
3	2013-06-07	NaT	389 days	NaT
4	2014-08-22	NaT	1040 days	NaT

Exploratory data analysis. Look at the distribution of seniority.

```
In [308]: import seaborn as sns
```

```
sns.distplot(Empl_reten['seniority'], hist = True, kde = True, bins = 30,
              kde_kws = {'shade': True, 'linewidth': 3})
```

```
Out[308]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1d4bc978>
```



```
In [309]: Empl_reten['seniority'].describe()
```

```
Out[309]: count      24702.000000
          mean         14.127803
          std          8.089520
          min          1.000000
          25%          7.000000
          50%         14.000000
          75%         21.000000
          max         99.000000
          Name: seniority, dtype: float64
```

There appear to be some outliers - possibly data entry mistakes. It is highly unlikely that someone worked for 99 years prior to their new employment. Let's remove the outliers after calculating z-scores.

There are some outliers in this dataset - as it is highly unlikely someone worked for 99 years before getting hired for their current position. Need to get rid of the outliers.

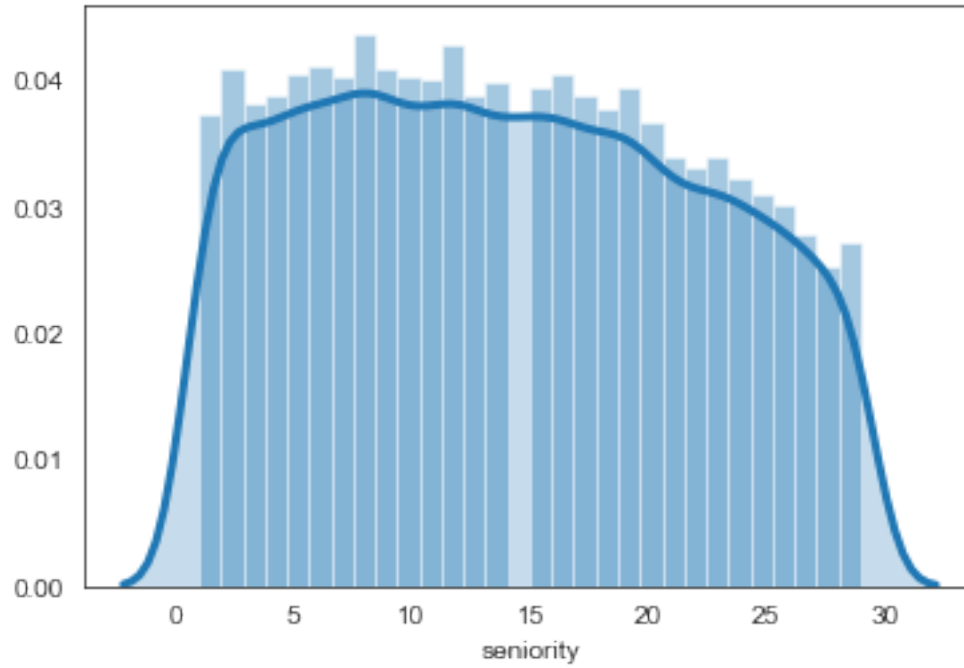
```
In [310]: from scipy import stats
```

```
z = np.abs(stats.zscore(Empl_reten['seniority']))
Empl_reten['seniority_z_score'] = z
Empl_reten = Empl_reten[Empl_reten['seniority_z_score'] < 4]
```

Seniority plot without outliers.

```
In [311]: sns.distplot(Empl_reten['seniority'], hist = True, kde = True, bins = 30,
                      kde_kws = {'shade': True, 'linewidth': 3})
```

```
Out[311]: <matplotlib.axes._subplots.AxesSubplot at 0x1a434a80b8>
```



Try removing rows with NaN values for employees who quit. Create a new data frame.

```
In [312]: empl_quit = Empl_reten.copy()
```

```
In [313]: empl_quit.head()
```

```
Out[313]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	

	quit_date	still_working	empl_quit_duration	still_working_duration	\
0	2015-10-30	NaT	585 days		NaT
1	2014-04-04	NaT	340 days		NaT
2	NaT	2015-12-13	NaT	426 days	
3	2013-06-07	NaT	389 days		NaT
4	2014-08-22	NaT	1040 days		NaT

```
seniority_z_score
```

0	1.714870
1	0.725916
2	0.015799
3	0.725916
4	1.096774

```
In [314]: empl_quit = empl_quit.drop(labels = ['still_working', 'still_working_duration'], axis=1)
```

```
In [315]: empl_quit.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24700 entries, 0 to 24699
Data columns (total 9 columns):
employee_id      24700 non-null float64
company_id       24700 non-null int64
dept             24700 non-null object
seniority         24700 non-null int64
salary           24700 non-null float64
join_date        24700 non-null datetime64[ns]
quit_date        13508 non-null datetime64[ns]
empl_quit_duration 13508 non-null timedelta64[ns]
seniority_z_score 24700 non-null float64
dtypes: datetime64[ns](2), float64(3), int64(2), object(1), timedelta64[ns](1)
memory usage: 1.9+ MB
```

```
In [316]: empl_quit_no_na = empl_quit.dropna()
```

```
In [317]: empl_quit_no_na.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 13508 entries, 0 to 24699
Data columns (total 9 columns):
employee_id      13508 non-null float64
company_id       13508 non-null int64
dept             13508 non-null object
seniority         13508 non-null int64
salary           13508 non-null float64
join_date        13508 non-null datetime64[ns]
quit_date        13508 non-null datetime64[ns]
empl_quit_duration 13508 non-null timedelta64[ns]
seniority_z_score 13508 non-null float64
dtypes: datetime64[ns](2), float64(3), int64(2), object(1), timedelta64[ns](1)
memory usage: 1.0+ MB
```

```
In [325]: empl_quit_no_na['empl_quit_days'] = empl_quit_no_na['empl_quit_duration'].astype('timedelta64[D]')
empl_quit_no_na['empl_quit_days'] = empl_quit_no_na['empl_quit_days'].astype(int)
```

```
/Users/Georgiy/anaconda3/envs/insight/lib/python3.7/site-packages/ipykernel_launcher.py:1: Set
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.htm>

```
"""Entry point for launching an IPython kernel.
```

```
/Users/Georgiy/anaconda3/envs/insight/lib/python3.7/site-packages/ipykernel_launcher.py:2: Set
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

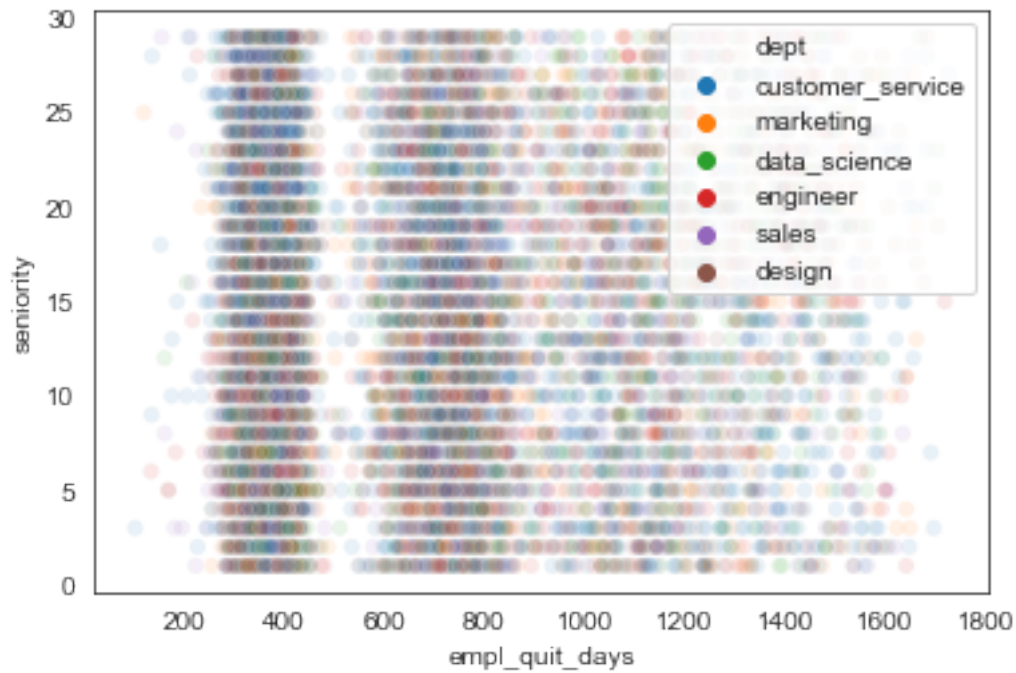
See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.htm>

```
In [326]: plt.scatter(empl_quit_no_na['empl_quit_days'], empl_quit_no_na['seniority'],
                    color = 'dodgerblue', alpha = 0.04)
plt.show() # Depending on whether you use IPython or interactive mode, etc.
```



```
In [327]: sns.scatterplot(empl_quit_no_na['empl_quit_days'],
                        empl_quit_no_na['seniority'], hue = empl_quit_no_na['dept'], alpha =
```


```
Out[327]: <matplotlib.axes._subplots.AxesSubplot at 0x1a42d25a20>
```



```
In [156]: sns.distplot(empl_quit_no_na['empl_quit_days'], hist = True, kde = True,
                    kde_kws = {'shade': True, 'linewidth': 3})
```

```
/Users/Georgiy/anaconda3/envs/insight/lib/python3.7/site-packages/scipy/stats/stats.py:1713: F
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[156]: <matplotlib.axes._subplots.AxesSubplot at 0x1a191f0f60>
```



It seems like there are three groups - maybe four. One group tends to leave after about a year (~375 days), a second group leaves after about 2 years (750 days), and a third group works for roughly 3 years (maybe ~1100 days). Then there is a fourth group that leaves maybe after 4 years or so (~1400 days).

Try removing rows with NaN values for employees who didn't quit. Create a new data frame.

```
In [328]: empl_still_working = Empl_reten.copy()
```

```
In [329]: empl_still_working = empl_still_working.drop(labels = ['quit_date', 'empl_quit_durat
```

```
In [330]: empl_still_working.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 24700 entries, 0 to 24699
```

```
Data columns (total 9 columns):
```

```
employee_id      24700 non-null float64
```

```
company_id       24700 non-null int64
```

```
dept             24700 non-null object
```

```
seniority        24700 non-null int64
```

```
salary          24700 non-null float64
```

```
join_date        24700 non-null datetime64[ns]
```

```
still_working    11192 non-null datetime64[ns]
```

```
still_working_duration 11192 non-null timedelta64[ns]
```

```
seniority_z_score 24700 non-null float64
```

```
dtypes: datetime64[ns](2), float64(3), int64(2), object(1), timedelta64[ns](1)
```

```
memory usage: 1.9+ MB
```



```
In [331]: empl_still_working_no_na = empl_still_working.dropna()
```

```
In [332]: empl_still_working_no_na.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 11192 entries, 2 to 24698
Data columns (total 9 columns):
employee_id      11192 non-null float64
company_id       11192 non-null int64
dept             11192 non-null object
seniority         11192 non-null int64
salary           11192 non-null float64
join_date        11192 non-null datetime64[ns]
still_working     11192 non-null datetime64[ns]
still_working_duration 11192 non-null timedelta64[ns]
seniority_z_score 11192 non-null float64
dtypes: datetime64[ns](2), float64(3), int64(2), object(1), timedelta64[ns](1)
memory usage: 874.4+ KB
```

```
In [333]: empl_still_working_no_na['empl_still_working_days'] = empl_still_working_no_na['still_working']
```

```
/Users/Georgiy/anaconda3/envs/insight/lib/python3.7/site-packages/ipykernel_launcher.py:1: SetitemError:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html
    """Entry point for launching an IPython kernel.
```

```
In [334]: empl_still_working_no_na['empl_still_working_days'] = empl_still_working_no_na['empl_still_working']
```

```
/Users/Georgiy/anaconda3/envs/insight/lib/python3.7/site-packages/ipykernel_launcher.py:1: SetitemError:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html
    """Entry point for launching an IPython kernel.
```

```
In [335]: empl_still_working_no_na.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 11192 entries, 2 to 24698
Data columns (total 10 columns):
employee_id      11192 non-null float64
company_id       11192 non-null int64
dept             11192 non-null object
seniority         11192 non-null int64
```

```


salary                11192 non-null float64
join_date              11192 non-null datetime64[ns]
still_working          11192 non-null datetime64[ns]
still_working_duration 11192 non-null timedelta64[ns]
seniority_z_score      11192 non-null float64
empl_still_working_days 11192 non-null int64
dtypes: datetime64[ns](2), float64(3), int64(3), object(1), timedelta64[ns](1)
memory usage: 961.8+ KB

```

```

In [336]: plt.scatter(empl_still_working_no_na['empl_still_working_days'], empl_still_working_no_na['seniority_z_score'],
                    color = 'dodgerblue', alpha = 0.04)
plt.show() # Depending on whether you use IPython or interactive mode, etc.

```



```


In [337]: sns.scatterplot(empl_still_working_no_na['empl_still_working_days'],
                        empl_still_working_no_na['seniority'], hue = empl_still_working_no_na['seniority'])

```

```

Out[337]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1d48edd8>


```



```
In [338]: sns.distplot(empl_still_working_no_na['empl_still_working_days'], hist = True, kde =
           kde_kws = {'shade': True, 'linewidth': 3})
```

```
/Users/Georgiy/anaconda3/envs/insight/lib/python3.7/site-packages/scipy/stats/stats.py:1713: F
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[338]: <matplotlib.axes._subplots.AxesSubplot at 0x1a43cf0e80>
```



Look at the number of employees withing each department (customer_service, engineering, etc.)


```
In [339]: Empl_reten['dept'].value_counts()
```

```
Out[339]: customer_service    9180
          engineer           4612
          data_science       3190
          sales              3172
          marketing          3166
          design             1380
          Name: dept, dtype: int64
```

```
In [340]: sns.set_style("white")
          sns.barplot(x = Empl_reten['dept'].value_counts().index,
                      y = Empl_reten['dept'].value_counts().values,
                      palette='muted', errcolor='dimgray')

          plt.xticks(rotation=90) # rotates the labels on the x-axis.
```

```
Out[340]: (array([0, 1, 2, 3, 4, 5]), <a list of 6 Text xticklabel objects>)
```



Sanity check: see if there are any missing values (NaN) in the data frame that includes employees who're still working.

```
In [341]: empl_still_working_no_na.isna().sum()
```

```
Out[341]: employee_id      0
          company_id      0
          dept            0
          seniority       0
          salary          0
          join_date       0
          still_working    0
          still_working_duration  0
          seniority_z_score  0
          empl_still_working_days  0
          dtype: int64
```


Sanity check: see if there are any missing values (NaN) in the data frame that includes employees who quit.

```
In [161]: empl_quit_no_na.isna().sum()
```

```
Out[161]: employee_id      0
          company_id      0
          dept             0
          seniority        0
          salary           0
          join_date        0
          quit_date        0
          empl_quit_duration 0
          seniority_z_score 0
          empl_quit_days    0
          dtype: int64
```


```
In [342]: sns.distplot(Empl_reten_no_na_no_out['salary'], hist = True, kde = True,
                      kde_kws = {'shade': True, 'linewidth': 3})
```

```
Out[342]: <matplotlib.axes._subplots.AxesSubplot at 0x1a442a3f28>
```



```
In [343]: sns.scatterplot(empl_quit_no_na['empl_quit_days'],
                          empl_quit_no_na['salary'], hue = empl_quit_no_na['dept'], alpha = 0.)
```

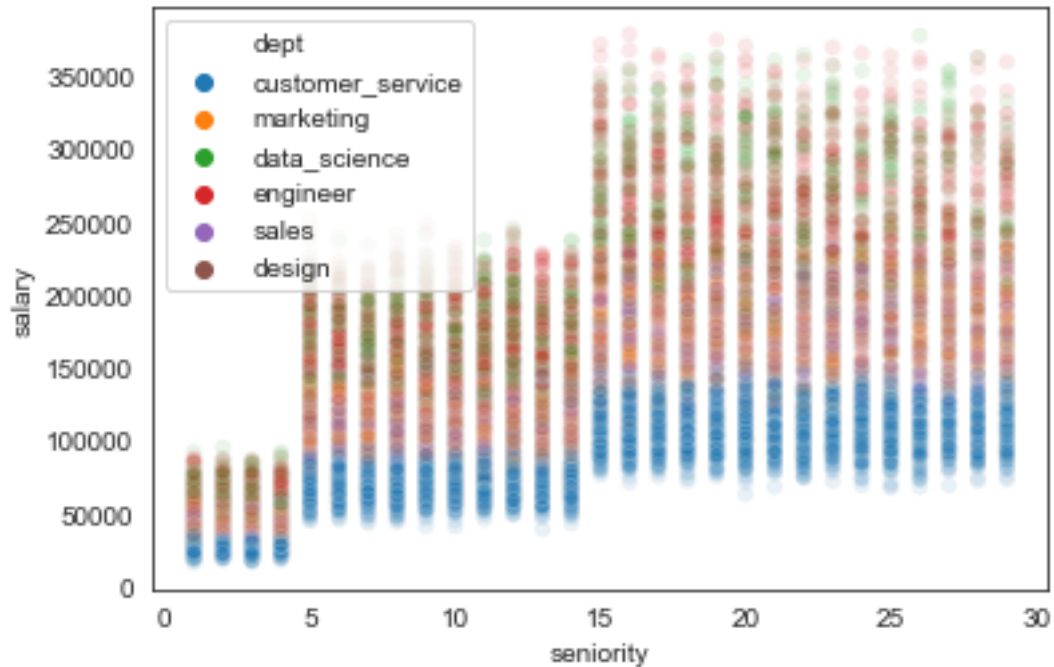
```
Out[343]: <matplotlib.axes._subplots.AxesSubplot at 0x1a443c3a20>
```



It looks like the average salary may stay to be about the same regardless of how long you've worked at a company. Break down the employment duration in categories, and see if the means for the categories are different.

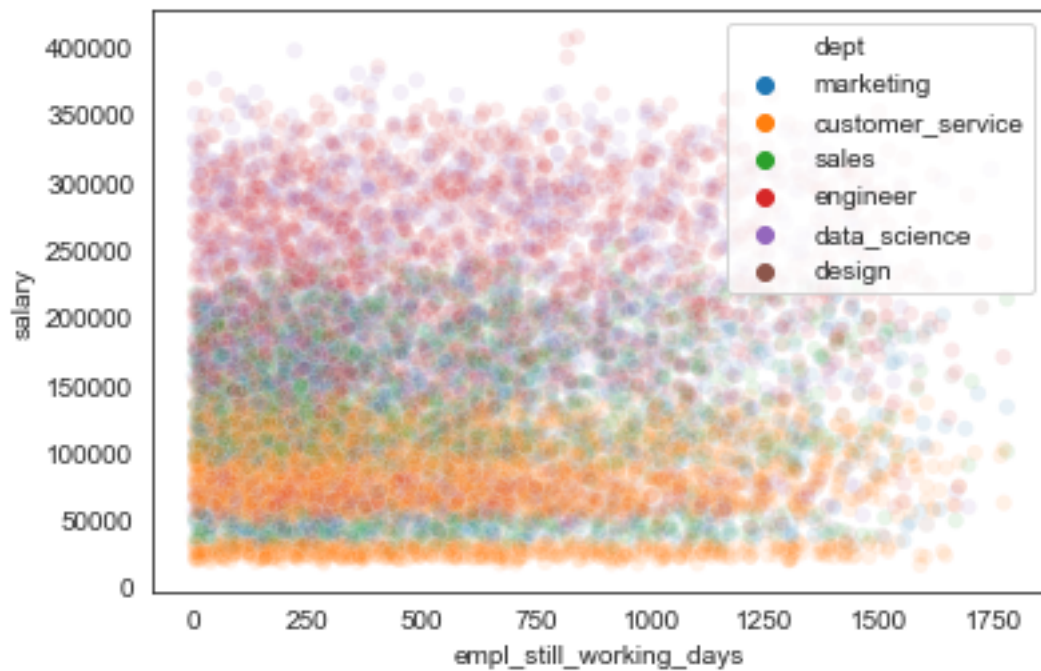
```
In [344]: sns.scatterplot(empl_quit_no_na['seniority'],  
                           empl_quit_no_na['salary'], hue = empl_quit_no_na['dept'], alpha = 0.1)
```

```
Out[344]: <matplotlib.axes._subplots.AxesSubplot at 0x1a44611240>
```




```
In [345]: sns.scatterplot(empl_still_working_no_na['empl_still_working_days'],
                        empl_still_working_no_na['salary'], hue = empl_still_working_no_na['dept'])
```

```
Out[345]: <matplotlib.axes._subplots.AxesSubplot at 0x1a4465aeb8>
```





```
In [346]: sns.scatterplot(empl_still_working_no_na['salary'],  
                           empl_still_working_no_na['empl_still_working_days'], alpha = 0.1)
```

```
Out[346]: <matplotlib.axes._subplots.AxesSubplot at 0x1a448b5a58>
```



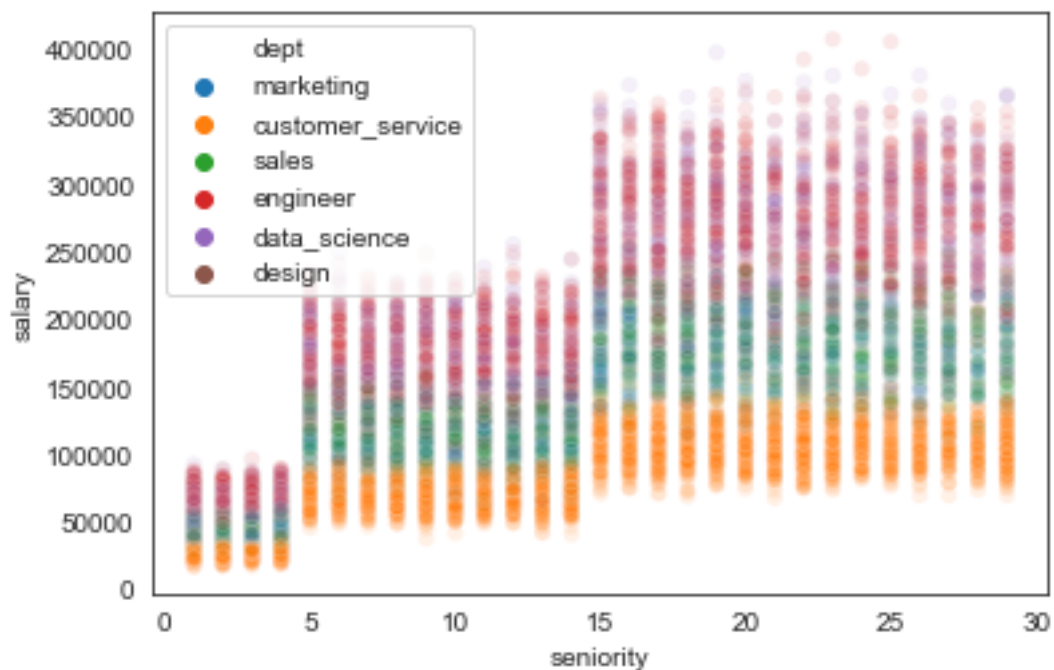
```
In [347]: sns.scatterplot(empl_quit_no_na['salary'],  
                           empl_quit_no_na['empl_quit_days'], alpha = 0.1)
```

```
Out[347]: <matplotlib.axes._subplots.AxesSubplot at 0x1a44a65b70>
```



```
In [348]: sns.scatterplot(empl_still_working_no_na['seniority'],
                           empl_still_working_no_na['salary'], hue = empl_still_working_no_na['dept'])
```

```
Out[348]: <matplotlib.axes._subplots.AxesSubplot at 0x1a44483c88>
```



Let's look at any linear relationships between salary and duration of employment. Do salary increases increase employment duration?

```
In [349]: empl_quit_SLR_df = empl_quit_no_na[['salary', 'empl_quit_days']]
```

```
In [350]: X = empl_quit_SLR_df['salary']
          y = empl_quit_SLR_df['empl_quit_days']
```

```
X = X.values
y = y.values
```

```
In [351]: from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
X_train = X_train.reshape(-1, 1)
X_test = X_test.reshape(-1, 1)
```

```
In [352]: from sklearn.linear_model import LinearRegression
          Model_1 = LinearRegression()
          Model_1.fit(X_train, y_train)
```

```
Out[352]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                           normalize=False)
```

```
In [353]: y_pred_1 = Model_1.predict(X_test)
```

```
In [354]: from sklearn.metrics import r2_score
```

```
In [355]: r2 = r2_score(y_test, y_pred_1)
          r2
```

```
Out[355]: 0.006936262289984452
```

```
In [356]: empl_still_working_SLR_df = empl_still_working_no_na[['salary', 'empl_still_working_days']]
```

```
In [357]: X_2 = empl_still_working_SLR_df['salary']
          y_2 = empl_still_working_SLR_df['empl_still_working_days']
```

```
X_2 = X_2.values
y_2 = y_2.values
```

```
In [358]: X_2_train, X_2_test, y_2_train, y_2_test = train_test_split(X_2, y_2, test_size=0.2, random_state=42)
```

```
X_2_train = X_2_train.reshape(-1, 1)
X_2_test = X_2_test.reshape(-1, 1)
```

```
In [359]: Model_2 = LinearRegression()
          Model_2.fit(X_2_train, y_2_train)
```

```
Out [359]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                             normalize=False)
```

```
In [360]: y_pred_2 = Model_2.predict(X_2_test)
```

```
In [361]: r2_2 = r2_score(y_2_test, y_pred_2)
          r2_2
```

```
Out [361]: 0.001342894018013463
```

Need to normalize response variables - that could explain poor model fit.

```
In [362]: empl_quit_no_na['log_salary'] = empl_quit_no_na['salary'].apply(np.log)
```

```
/Users/Georgiy/anaconda3/envs/insight/lib/python3.7/site-packages/ipykernel_launcher.py:1: Set
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
"""Entry point for launching an IPython kernel.

```
In [363]: empl_still_working_no_na['log_salary'] = empl_still_working_no_na['salary'].apply(np
```


```
/Users/Georgiy/anaconda3/envs/insight/lib/python3.7/site-packages/ipykernel_launcher.py:1: Set
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
"""Entry point for launching an IPython kernel.

```
In [364]: sns.distplot(empl_quit_no_na['log_salary'], hist = True, kde = True,
                      kde_kws = {'shade': True, 'linewidth': 3})
```


```
/Users/Georgiy/anaconda3/envs/insight/lib/python3.7/site-packages/scipy/stats/stats.py:1713: F
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out [364]: <matplotlib.axes._subplots.AxesSubplot at 0x1a44c48358>
```



```
In [365]: sns.distplot(empl_still_working_no_na['log_salary'], hist = True, kde = True,  
                      kde_kws = {'shade': True, 'linewidth': 3})
```

```
Out[365]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1b4edc18>
```



```
In [366]: Empl_reten.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24700 entries, 0 to 24699
Data columns (total 12 columns):
employee_id      24700 non-null float64
company_id       24700 non-null int64
dept             24700 non-null object
seniority        24700 non-null int64
salary           24700 non-null float64
join_date        24700 non-null datetime64[ns]
quit_date        13508 non-null datetime64[ns]
still_working    11192 non-null datetime64[ns]
empl_quit_duration 13508 non-null timedelta64[ns]
still_working_duration 11192 non-null timedelta64[ns]
seniority_z_score 24700 non-null float64
empl_quit_days   13508 non-null float64
dtypes: datetime64[ns](3), float64(4), int64(2), object(1), timedelta64[ns](2)
memory usage: 2.4+ MB
```

```
In [367]: Empl_reten['empl_still_working_days'] = Empl_reten['still_working_duration'].astype('D')
```

```
In [368]: Empl_reten['empl_still_working_days'] = Empl_reten['empl_still_working_days'].fillna(0)
```

```
In [369]: Empl_reten.head(10)
```

```
Out[369]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	
5	509529.0	4	data_science	14	165000.0	2012-01-30	
6	88600.0	4	customer_service	21	107000.0	2013-10-21	
7	716309.0	2	customer_service	4	30000.0	2014-03-05	
8	172999.0	9	engineer	7	160000.0	2012-12-10	
9	504159.0	1	sales	7	104000.0	2012-06-12	

	quit_date	still_working	empl_quit_duration	still_working_duration	\
0	2015-10-30	NaT	585 days	NaT	
1	2014-04-04	NaT	340 days	NaT	
2	NaT	2015-12-13	NaT	426 days	
3	2013-06-07	NaT	389 days	NaT	
4	2014-08-22	NaT	1040 days	NaT	
5	2013-08-30	NaT	578 days	NaT	

6	NaT	2015-12-13	NaT	783 days
7	NaT	2015-12-13	NaT	648 days
8	2015-10-23	NaT	1047 days	NaT
9	NaT	2015-12-13	NaT	1279 days

	seniority_z_score	empl_quit_days	empl_still_working_days
0	1.714870	585.0	False
1	0.725916	340.0	False
2	0.015799	NaN	426
3	0.725916	389.0	False
4	1.096774	1040.0	False
5	0.015799	578.0	False
6	0.849536	NaN	783
7	1.251991	NaN	648
8	0.881134	1047.0	False
9	0.881134	NaN	1279

```
In [370]: Empl_reten['quit_date'] = Empl_reten['quit_date'].fillna(value = False)
```

```
In [371]: Empl_reten['still_working'] = Empl_reten['still_working'].fillna(value = False)
```

```
In [372]: Empl_reten['empl_quit_duration'] = Empl_reten['empl_quit_duration'].fillna(value = False)
```

```
In [373]: Empl_reten['still_working_duration'] = Empl_reten['still_working_duration'].fillna(value = False)
```

```
In [374]: Empl_reten['empl_quit_days'] = Empl_reten['empl_quit_days'].fillna(value = False)
```

```
In [375]: Empl_reten.head(10)
```

```
Out[375]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	
5	509529.0	4	data_science	14	165000.0	2012-01-30	
6	88600.0	4	customer_service	21	107000.0	2013-10-21	
7	716309.0	2	customer_service	4	30000.0	2014-03-05	
8	172999.0	9	engineer	7	160000.0	2012-12-10	
9	504159.0	1	sales	7	104000.0	2012-06-12	

	quit_date	still_working	empl_quit_duration	\
0	2015-10-30 00:00:00	False	585 days 00:00:00	
1	2014-04-04 00:00:00	False	340 days 00:00:00	
2	False	2015-12-13 00:00:00	False	
3	2013-06-07 00:00:00	False	389 days 00:00:00	
4	2014-08-22 00:00:00	False	1040 days 00:00:00	
5	2013-08-30 00:00:00	False	578 days 00:00:00	
6	False	2015-12-13 00:00:00	False	

7		False	2015-12-13 00:00:00	False
8	2015-10-23 00:00:00		False	1047 days 00:00:00
9		False	2015-12-13 00:00:00	False

	still_working_duration	seniority_z_score	empl_quit_days \
0	False	1.714870	585
1	False	0.725916	340
2	426 days 00:00:00	0.015799	False
3	False	0.725916	389
4	False	1.096774	1040
5	False	0.015799	578
6	783 days 00:00:00	0.849536	False
7	648 days 00:00:00	1.251991	False
8	False	0.881134	1047
9	1279 days 00:00:00	0.881134	False

	empl_still_working_days
0	False
1	False
2	426
3	False
4	False
5	False
6	783
7	648
8	False
9	1279

In [376]: Empl_reten = Empl_reten.drop(labels = ['empl_quit_duration', 'still_working_duration'])

In [377]: Empl_reten.head(10)

Out[377]:

	employee_id	company_id	dept	seniority	salary	join_date \
0	13021.0	7	customer_service	28	89000.0	2014-03-24
1	825355.0	7	marketing	20	183000.0	2013-04-29
2	927315.0	4	marketing	14	101000.0	2014-10-13
3	662910.0	7	customer_service	20	115000.0	2012-05-14
4	256971.0	2	data_science	23	276000.0	2011-10-17
5	509529.0	4	data_science	14	165000.0	2012-01-30
6	88600.0	4	customer_service	21	107000.0	2013-10-21
7	716309.0	2	customer_service	4	30000.0	2014-03-05
8	172999.0	9	engineer	7	160000.0	2012-12-10
9	504159.0	1	sales	7	104000.0	2012-06-12

	quit_date	still_working	seniority_z_score	empl_quit_days \
0	2015-10-30 00:00:00	False	1.714870	585
1	2014-04-04 00:00:00	False	0.725916	340
2	False	2015-12-13 00:00:00	0.015799	False

3	2013-06-07 00:00:00	False	0.725916	389
4	2014-08-22 00:00:00	False	1.096774	1040
5	2013-08-30 00:00:00	False	0.015799	578
6	False 2015-12-13 00:00:00		0.849536	False
7	False 2015-12-13 00:00:00		1.251991	False
8	2015-10-23 00:00:00	False	0.881134	1047
9	False 2015-12-13 00:00:00		0.881134	False

	empl_still_working_days
0	False
1	False
2	426
3	False
4	False
5	False
6	783
7	648
8	False
9	1279

```
In [378]: Empl_reten['Class_quit'] = Empl_reten['empl_still_working_days'] == False
```

```
In [379]: Empl_reten['Class_quit'] = Empl_reten['Class_quit'].astype(int)
```

Create dummy variables for different departments

```
In [380]: Empl_reten = pd.get_dummies(Empl_reten, prefix = ['dept_dummy'], columns = ['dept'],
```

```
In [381]: Empl_reten.head())
```

```
Out[381]:
```

	employee_id	company_id	seniority	salary	join_date	\
0	13021.0	7	28	89000.0	2014-03-24	
1	825355.0	7	20	183000.0	2013-04-29	
2	927315.0	4	14	101000.0	2014-10-13	
3	662910.0	7	20	115000.0	2012-05-14	
4	256971.0	2	23	276000.0	2011-10-17	

	quit_date	still_working	seniority_z_score	empl_quit_days	\
0	2015-10-30 00:00:00	False	1.714870	585	
1	2014-04-04 00:00:00	False	0.725916	340	
2	False 2015-12-13 00:00:00		0.015799	False	
3	2013-06-07 00:00:00	False	0.725916	389	
4	2014-08-22 00:00:00	False	1.096774	1040	

	empl_still_working_days	Class_quit	dept_dummy_data_science	\
0	False	1	0	
1	False	1	0	
2	426	0	0	
3	False	1	0	

	False	1	1
dept_dummy_design			
0	0	0	0
1	0	0	1
2	0	0	1
3	0	0	0
4	0	0	0

	dept_dummy_sales
0	0
1	0
2	0
3	0
4	0

```
In [391]: X_rf = Empl_reten[['seniority', 'salary', 'company_id', 'dept_dummy_data_science', 'dept_dummy_engineer', 'dept_dummy_marketing', 'dept_dummy_sales']]
y_rf = Empl_reten['Class_quit']
```

```
In [392]: from sklearn.ensemble import RandomForestClassifier
```

```
In [393]: rf = RandomForestClassifier(n_estimators = 1000, random_state = 42)
```

```
In [394]: X_rf_train, X_rf_test, y_rf_train, y_rf_test = train_test_split(X_rf, y_rf, test_size=0.2)
```

```
In [395]: rf.fit(X_rf_train, y_rf_train)
```

```
Out[395]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=1000, n_jobs=None,
oob_score=False, random_state=42, verbose=0, warm_start=False)
```

```
In [396]: predictions_rf = rf.predict(X_rf_test)
```

```
In [397]: errors = abs(predictions_rf - y_rf_test)
```

```
In [398]: print('Mean Absolute Error:', round(np.mean(errors), 2))
```

Mean Absolute Error: 0.5

```
In [399]: # Calculate mean absolute percentage error (MAPE)
mape = 100 * (errors / y_rf_test)
# Calculate and display accuracy
accuracy = 100 - np.mean(mape)
print('Accuracy:', round(accuracy, 2), '%.')
```

