

Steph_Gervasi_Employee_Retention

February 20, 2019

```
In [2]: # Import the basics for some initial EDA
```

```
import numpy as np
import pandas as pd
from scipy import stats
import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0)
import seaborn as sns
```

```
In [3]: # Some information I've been given about the data:
```

```
##### Columns I have:
## employee_id : id of the employee. Unique by employee per company
## company_id : company id.
## dept : employee dept
## seniority : number of yrs of work experience when hired
## salary: avg yearly salary of the employee during her tenure within the company
## join_date: when the employee joined the company, it can only be between 2011/01/24
## quit_date: when the employee left her job (if she is still employed as of 2015/12/1)
```

```
In [4]: # Get the data
```

```
dc1 = pd.read_csv("employee_retention_data.csv")
```

```
In [5]: # Check data dimensions
```

```
print("The data has {0} rows and {1} columns".format(dc1.shape[0], dc1.shape[1]))
```

The data has 24702 rows and 7 columns

```
In [6]: # Take a quick look to see what the data look like
```

```
dc1.head()
```

```
Out[6]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	

```
4      256971.0      2      data_science      23  276000.0  2011-10-17
```

```
      quit_date
0  2015-10-30
1  2014-04-04
2           NaN
3  2013-06-07
4  2014-08-22
```

```
In [7]: # Confirm column names
        list(dc1.columns.values)
```

```
Out[7]: ['employee_id',
         'company_id',
         'dept',
         'seniority',
         'salary',
         'join_date',
         'quit_date']
```

```
In [8]: # Get info about data types and missingness
        dc1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24702 entries, 0 to 24701
Data columns (total 7 columns):
employee_id    24702 non-null float64
company_id     24702 non-null int64
dept           24702 non-null object
seniority      24702 non-null int64
salary         24702 non-null float64
join_date      24702 non-null object
quit_date      13510 non-null object
dtypes: float64(2), int64(2), object(3)
memory usage: 1.3+ MB
```

```
In [9]: # Check percentage of missing values in the columns to decide whether to drop, impute,
        miss = dc1.isnull().sum()/len(dc1)
        print(type(miss))
        miss = miss[miss > 0]
        miss.sort_values(inplace = True)
        miss
        # I know that these data are missing because it means that the people still work there
```

```
<class 'pandas.core.series.Series'>
```


```
Out[9]: quit_date    0.453081
        dtype: float64
```

```
In [10]: # Now look at the other variables/columns in the dataset
# First look at dept variable
dc1['dept'].unique()
# there are 6
```

```
Out[10]: array(['customer_service', 'marketing', 'data_science', 'engineer',
               'sales', 'design'], dtype=object)
```

```
In [11]: # What is the break down/count of number of employees by department across all companies
sns.set(style="whitegrid", color_codes=True)
plt.rcParams['figure.figsize'] = (10.0, 8.0)
sns.set_context("paper", rc={"font.size":16,"axes.titlesize":16,"axes.labelsize":16,"figure.figsize":(10,8)})
ax = sns.countplot(y="dept", data=dc1)
ax.set(xlabel='Overall Employee Dept Affiliation', ylabel='Count')
plt.show()

# CONCLUDE:
# Greatest bulk of employees are coming from customer service category. May want to c
```




```
In [206]: # What is the break down/count of number of employees by department shown for each company
sns.set(style="whitegrid", color_codes=True)
plt.rcParams['figure.figsize'] = (10.0, 10.0)
sns.set_context("paper", rc={"font.size":16,"axes.titlesize":16,"axes.labelsize":16,"figure.figsize":(10,10)})
```

```
ax = sns.countplot(y="company_id", hue = "dept", data=dc1)
ax.set(xlabel='Overall Employee Dept Affiliation', ylabel='Count')
plt.show()
```

CONCLUDE:

Trend observed above seems to also be present in the within company context : most

Company 11 and 12 are very different from the rest. May want to think about includ



In [13]: *# Next look at the seniority variable:*

```
np.sort(dc1['seniority'].unique())
```

It is highly unlikely that someone has a seniority status of 98 or np.sort(dc1['seniority'].unique())


Out[13]: array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 98, 99])

```
In [14]: # Quickly, look at seniority and salary relationship (see both distribution/histogram
sns.jointplot(x="seniority", y="salary", data=dc1)
```

```
# CONCLUDE:
```


```
# Need to get rid of those outliers - don't have any information on how to correct th
```

```
Out[14]: <seaborn.axisgrid.JointGrid at 0x102bc1588>
```



```
In [15]: # plot distribution of seniority values - outliers evident here too
sns.distplot(dc1['seniority'], rug = True, color = 'y')
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x1a09baed30>
```



```
In [16]: # deal with outliers in seniority column, here only 2, so will remove. Check shape before
         dc1.shape
```

```
Out[16]: (24702, 7)
```


```
In [17]: # Process of removing seniority outliers based on being above 3 st deviations of the mean
         target = dc1['seniority']
         mean = target.mean()
         sd = target.std()
         reduced_dc1 = dc1[(target > mean - 3*sd) & (target < mean + 3*sd)]
         reduced_dc1.shape
```

```
# this removed 2 outliers that were beyond 3 sd of the mean.
```

```
Out[17]: (24700, 7)
```

```
In [18]: # Now take a look at the variable, salary, by itself
         sns.distplot(reduced_dc1['salary'], rug = True, color = 'y')
         # looks like some outliers here, too, at the high end. These could interfere with pre
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x1a15240b00>
```



```
In [19]: ## Here applying the same process/criteria for outlier removal in the salary column a
target2 = reduced_dc1['salary']
mean2 = target2.mean()
sd2 = target2.std()
reduced2_dc1 = reduced_dc1[(target2 > mean2 - 3*sd2) & (target2 < mean2 + 3*sd2)]
reduced2_dc1.shape
# As a result, this removed 22 additional data points.
```

```
Out[19]: (24678, 7)
```

```
In [20]: # give the new dataframe a more intuitive name after outlier removal from the seniori
dc1_NEW = reduced2_dc1
```

```
In [21]: # Our reduced dataframe:
dc1_NEW.head()
```

```
Out[21]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	

```

4      256971.0      2      data_science      23  276000.0  2011-10-17

      quit_date
0  2015-10-30
1  2014-04-04
2      NaN
3  2013-06-07
4  2014-08-22

```

```

In [22]: # Now, look at company_id variable
dc1_NEW['company_id'].nunique()
# size of companies clearly varies 11 and 12 are smaller than the rest - this could b

```

```

Out[22]: 12

```

```

In [23]: # Now, check the employee_id variable (make sure no duplicate IDs; # rows shoudl = #
dc1_NEW['employee_id'].nunique()

```

```

Out[23]: 24678

```

```

In [ ]: # The next steps will be to join and quit dates: first take a look and then convert th

```

```

In [236]: # First do a quick visual scan to make sure they're all in the same format (although
#dc1_NEW['join_date'].value_counts(dropna = False)

```

```

In [237]: #dc1_NEW['quit_date'].value_counts(dropna = False)

```

```

In [26]: # get all dates in standard format before calculating difference in time to get tenure
import datetime
dc1_NEW['join_date'] = pd.to_datetime(dc1_NEW['join_date'])
dc1_NEW['quit_date'] = pd.to_datetime(dc1_NEW['quit_date'])
dc1_NEW.head()

```

```

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
This is separate from the ipykernel package so we can avoid doing imports until

```

Out[26]:   employee_id  company_id      dept  seniority  salary  join_date  \
0      13021.0      7  customer_service      28  89000.0  2014-03-24

```


1	825355.0	7	marketing	20	183000.0	2013-04-29
2	927315.0	4	marketing	14	101000.0	2014-10-13
3	662910.0	7	customer_service	20	115000.0	2012-05-14
4	256971.0	2	data_science	23	276000.0	2011-10-17

```

quit_date
0 2015-10-30
1 2014-04-04
2          NaT
3 2013-06-07
4 2014-08-22

```

```
In [27]: # Check updated format of date columns
dc1_NEW.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 24678 entries, 0 to 24699
Data columns (total 7 columns):
employee_id    24678 non-null float64
company_id     24678 non-null int64
dept           24678 non-null object
seniority       24678 non-null int64
salary         24678 non-null float64
join_date      24678 non-null datetime64[ns]
quit_date      13501 non-null datetime64[ns]
dtypes: datetime64[ns](2), float64(2), int64(2), object(1)
memory usage: 1.5+ MB

```

```
In [29]: # Calcualte a new column for total tenure, which is the time someone has worked at the company
dc1_NEW['total_tenure'] = dc1_NEW['quit_date'] - dc1_NEW['join_date']
dc1_NEW.head()
# string and number are the output of this new column

```

```

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```

"""Entry point for launching an IPython kernel.

```

```

Out[29]:
employee_id  company_id  dept  seniority  salary  join_date  \
0      13021.0         7  customer_service    28   89000.0  2014-03-24
1      825355.0         7      marketing    20  183000.0  2013-04-29
2      927315.0         4      marketing    14  101000.0  2014-10-13
3      662910.0         7  customer_service    20  115000.0  2012-05-14
4      256971.0         2      data_science    23  276000.0  2011-10-17

```

	quit_date	total_tenure
0	2015-10-30	585 days
1	2014-04-04	340 days
2	NaT	NaT
3	2013-06-07	389 days
4	2014-08-22	1040 days

```
In [30]: # Convert the variable total_tenure that was just created so that it only includes the
dc1_NEW['days_employeed'] = (dc1_NEW['quit_date'] - dc1_NEW['join_date']).dt.days
dc1_NEW.head()
# just number of days and NaNs in this new column, days_employeed
```

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
 """Entry point for launching an IPython kernel.

```
Out[30]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	

	quit_date	total_tenure	days_employeed
0	2015-10-30	585 days	585.0
1	2014-04-04	340 days	340.0
2	NaT	NaT	NaN
3	2013-06-07	389 days	389.0
4	2014-08-22	1040 days	1040.0

```
In [32]: # Change the NAs in the new days_employeed column to something meaningful, because the
dc1_NEW['days_employeed'].fillna('currently_employeed', inplace = True)
dc1_NEW.head()
```

/anaconda3/lib/python3.6/site-packages/pandas/core/generic.py:5434: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
 self._update_inplace(new_data)

```
Out[32]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	

3	662910.0	7	customer_service	20	115000.0	2012-05-14
4	256971.0	2	data_science	23	276000.0	2011-10-17

	quit_date	total_tenure	days_employeed
0	2015-10-30	585 days	585
1	2014-04-04	340 days	340
2	NaT	NaT	currently_employeed
3	2013-06-07	389 days	389
4	2014-08-22	1040 days	1040

```
In [33]: # Also create new columns for start year and month since this could potentially tell
dc1_NEW['start_year'] = (dc1_NEW['join_date'].dt.year)
dc1_NEW['start_month'] = (dc1_NEW['join_date'].dt.month)
dc1_NEW.head()
```

```
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
This is separate from the ipykernel package so we can avoid doing imports until

```
Out[33]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	

	quit_date	total_tenure	days_employeed	start_year	start_month
0	2015-10-30	585 days	585	2014	3
1	2014-04-04	340 days	340	2013	4
2	NaT	NaT	currently_employeed	2014	10
3	2013-06-07	389 days	389	2012	5
4	2014-08-22	1040 days	1040	2011	10

```
In [ ]: # May also want to take the continuous column/feature called days_employeed and turn i
```

```
In [34]: # First, duplicate the column and call it days_employed_cont
dc1_NEW['days_employeed_cont']=dc1_NEW['days_employeed']
```

```
# Now, turn the string 'curr_employeed' into a high number 1800 is 5 years so use 185
```

```
dc1_NEW['days_employeed_cont'] = dc1_NEW['days_employeed_cont'].replace('currently_employeed', 0)
dc1_NEW.head()
```

```
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
Out[34]:
```


	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	

	quit_date	total_tenure	days_employeed	start_year	start_month	\
0	2015-10-30	585 days	585	2014	3	
1	2014-04-04	340 days	340	2013	4	
2	NaT	NaT	currently_employeed	2014	10	
3	2013-06-07	389 days	389	2012	5	
4	2014-08-22	1040 days	1040	2011	10	

	days_employeed_cont
0	585.0
1	340.0
2	1850.0
3	389.0
4	1040.0

```
In [36]: # Now, the column days_employeed_cont is all numeric values, and we can inspect further
sns.distplot(dc1_NEW['days_employeed_cont'], kde=True, rug=False, norm_hist=False, color='b')
# looks like at least 4 major natural breakpoints where people worked for 1 year, ~ 2
```

```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x1a18b68c50>
```



```
In [37]: # After first trying to use the quartile method, it seems like it might be best to manually
        bins = [0, 360, 720, 1080, 1440, 1850]
        names = ['<1yr', '1-2yr', '2-3yr', '3-4yr', '4+yrs']

        # a new column called tenure_category is the binned times/durations that people were
        dc1_NEW['tenure_category'] = pd.cut(dc1_NEW['days_employed_cont'], bins, labels=names)
        dc1_NEW.head(20)

        #print(dc1_NEW.dtypes)

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html
"""
```

```
Out[37]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	

1	825355.0	7	marketing	20	183000.0	2013-04-29
2	927315.0	4	marketing	14	101000.0	2014-10-13
3	662910.0	7	customer_service	20	115000.0	2012-05-14
4	256971.0	2	data_science	23	276000.0	2011-10-17
5	509529.0	4	data_science	14	165000.0	2012-01-30
6	88600.0	4	customer_service	21	107000.0	2013-10-21
7	716309.0	2	customer_service	4	30000.0	2014-03-05
8	172999.0	9	engineer	7	160000.0	2012-12-10
9	504159.0	1	sales	7	104000.0	2012-06-12
10	892155.0	6	customer_service	13	72000.0	2012-11-12
11	904158.0	2	marketing	17	230000.0	2015-05-11
12	939058.0	1	marketing	1	48000.0	2012-12-10
13	163427.0	10	marketing	23	154000.0	2012-06-18
14	461248.0	2	sales	20	201000.0	2013-09-16
15	265226.0	1	data_science	4	80000.0	2014-05-27
16	932790.0	7	marketing	10	88000.0	2011-11-30
17	69693.0	7	customer_service	6	54000.0	2014-03-31
18	721600.0	2	marketing	20	193000.0	2014-12-29
19	982668.0	1	customer_service	14	76000.0	2015-07-27

	quit_date	total_tenure	days_employeed	start_year	start_month	\
0	2015-10-30	585 days	585	2014	3	
1	2014-04-04	340 days	340	2013	4	
2	NaT	NaT	currently_employeed	2014	10	
3	2013-06-07	389 days	389	2012	5	
4	2014-08-22	1040 days	1040	2011	10	
5	2013-08-30	578 days	578	2012	1	
6	NaT	NaT	currently_employeed	2013	10	
7	NaT	NaT	currently_employeed	2014	3	
8	2015-10-23	1047 days	1047	2012	12	
9	NaT	NaT	currently_employeed	2012	6	
10	2015-02-27	837 days	837	2012	11	
11	NaT	NaT	currently_employeed	2015	5	
12	2013-11-15	340 days	340	2012	12	
13	2015-09-25	1194 days	1194	2012	6	
14	2014-08-22	340 days	340	2013	9	
15	2015-07-10	409 days	409	2014	5	
16	2013-11-22	723 days	723	2011	11	
17	NaT	NaT	currently_employeed	2014	3	
18	NaT	NaT	currently_employeed	2014	12	
19	NaT	NaT	currently_employeed	2015	7	

	days_employeed_cont	tenure_category
0	585.0	1-2yr
1	340.0	<1yr
2	1850.0	4+yrs
3	389.0	1-2yr
4	1040.0	2-3yr

5	578.0	1-2yr
6	1850.0	4+yrs
7	1850.0	4+yrs
8	1047.0	2-3yr
9	1850.0	4+yrs
10	837.0	2-3yr
11	1850.0	4+yrs
12	340.0	<1yr
13	1194.0	3-4yr
14	340.0	<1yr
15	409.0	1-2yr
16	723.0	2-3yr
17	1850.0	4+yrs
18	1850.0	4+yrs
19	1850.0	4+yrs

```
In [42]: #Also, consider creating a lower resolution, binary column that tells us just whether
dc1_NEW['still_employeed'] = np.where(pd.isnull(dc1_NEW['quit_date']), '1', '0')
dc1_NEW.head(10)
```

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
Out[42]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	
5	509529.0	4	data_science	14	165000.0	2012-01-30	
6	88600.0	4	customer_service	21	107000.0	2013-10-21	
7	716309.0	2	customer_service	4	30000.0	2014-03-05	
8	172999.0	9	engineer	7	160000.0	2012-12-10	
9	504159.0	1	sales	7	104000.0	2012-06-12	

	quit_date	total_tenure	days_employeed	start_year	start_month	\
0	2015-10-30	585 days	585	2014	3	
1	2014-04-04	340 days	340	2013	4	
2	NaT	NaT	currently_employeed	2014	10	
3	2013-06-07	389 days	389	2012	5	
4	2014-08-22	1040 days	1040	2011	10	
5	2013-08-30	578 days	578	2012	1	
6	NaT	NaT	currently_employeed	2013	10	

7	NaT	NaT	currently_employeed	2014	3
8	2015-10-23	1047 days	1047	2012	12
9	NaT	NaT	currently_employeed	2012	6

	days_employeed_cont	tenure_category	still_employeed
0	585.0	1-2yr	0
1	340.0	<1yr	0
2	1850.0	4+yrs	1
3	389.0	1-2yr	0
4	1040.0	2-3yr	0
5	578.0	1-2yr	0
6	1850.0	4+yrs	1
7	1850.0	4+yrs	1
8	1047.0	2-3yr	0
9	1850.0	4+yrs	1

```
In [40]: #Look at counts of two created groups for still employeed or not to see if they are b
dc1_NEW['still_employeed'].value_counts(dropna = False)
```

```
Out[40]: 0    13501
         1    11177
         Name: still_employeed, dtype: int64
```

```
In [41]: # Sanity check
dc1_NEW.isnull().sum(axis = 0)
```

```
Out[41]: employee_id          0
         company_id          0
         dept                0
         seniority           0
         salary              0
         join_date           0
         quit_date          11177
         total_tenure        11177
         days_employeed       0
         start_year          0
         start_month         0
         days_employeed_cont  0
         tenure_category      0
         still_employeed      0
         dtype: int64
```

```
In [44]: # I also want to create a salary ranking column in case this is a better predictor th
dc1_NEW['salary_rank']=pd.qcut(dc1_NEW['salary'],4,labels=['low','med','high','very_h
dc1_NEW.head(20)
```

```
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
Out[44]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	
5	509529.0	4	data_science	14	165000.0	2012-01-30	
6	88600.0	4	customer_service	21	107000.0	2013-10-21	
7	716309.0	2	customer_service	4	30000.0	2014-03-05	
8	172999.0	9	engineer	7	160000.0	2012-12-10	
9	504159.0	1	sales	7	104000.0	2012-06-12	
10	892155.0	6	customer_service	13	72000.0	2012-11-12	
11	904158.0	2	marketing	17	230000.0	2015-05-11	
12	939058.0	1	marketing	1	48000.0	2012-12-10	
13	163427.0	10	marketing	23	154000.0	2012-06-18	
14	461248.0	2	sales	20	201000.0	2013-09-16	
15	265226.0	1	data_science	4	80000.0	2014-05-27	
16	932790.0	7	marketing	10	88000.0	2011-11-30	
17	69693.0	7	customer_service	6	54000.0	2014-03-31	
18	721600.0	2	marketing	20	193000.0	2014-12-29	
19	982668.0	1	customer_service	14	76000.0	2015-07-27	

	quit_date	total_tenure	days_employeed	start_year	start_month	\
0	2015-10-30	585 days	585	2014	3	
1	2014-04-04	340 days	340	2013	4	
2	NaT	NaT	currently_employeed	2014	10	
3	2013-06-07	389 days	389	2012	5	
4	2014-08-22	1040 days	1040	2011	10	
5	2013-08-30	578 days	578	2012	1	
6	NaT	NaT	currently_employeed	2013	10	
7	NaT	NaT	currently_employeed	2014	3	
8	2015-10-23	1047 days	1047	2012	12	
9	NaT	NaT	currently_employeed	2012	6	
10	2015-02-27	837 days	837	2012	11	
11	NaT	NaT	currently_employeed	2015	5	
12	2013-11-15	340 days	340	2012	12	
13	2015-09-25	1194 days	1194	2012	6	
14	2014-08-22	340 days	340	2013	9	
15	2015-07-10	409 days	409	2014	5	
16	2013-11-22	723 days	723	2011	11	
17	NaT	NaT	currently_employeed	2014	3	
18	NaT	NaT	currently_employeed	2014	12	
19	NaT	NaT	currently_employeed	2015	7	

	days_employed_cont	tenure_category	still_employed	salary_rank
0	585.0	1-2yr	0	med
1	340.0	<1yr	0	high
2	1850.0	4+yrs	1	med
3	389.0	1-2yr	0	med
4	1040.0	2-3yr	0	very_high
5	578.0	1-2yr	0	high
6	1850.0	4+yrs	1	med
7	1850.0	4+yrs	1	low
8	1047.0	2-3yr	0	high
9	1850.0	4+yrs	1	med
10	837.0	2-3yr	0	low
11	1850.0	4+yrs	1	very_high
12	340.0	<1yr	0	low
13	1194.0	3-4yr	0	high
14	340.0	<1yr	0	very_high
15	409.0	1-2yr	0	med
16	723.0	2-3yr	0	med
17	1850.0	4+yrs	1	low
18	1850.0	4+yrs	1	very_high
19	1850.0	4+yrs	1	low

In [70]: *# Similarly, I may want to just compare size of companies - could be important to take*
total_employees_count = dc1_NEW.groupby(['company_id']).size().reset_index(name='total_employees_count')
total_employees_count.head(12)

Out [70]:

	company_id	total_number_of_employees
0	1	8472
1	2	4213
2	3	2749
3	4	2062
4	5	1755
5	6	1291
6	7	1224
7	8	1047
8	9	961
9	10	864
10	11	16
11	12	24

In [72]: total_employees_by_dept = dc1_NEW.groupby(['company_id', 'dept']).size().reset_index(name='total_employees_by_dept')
total_employees_by_dept.tail(20) *# check tail for those small companies - 11 and 12.*

Out [72]:

	company_id	dept	number_of_employees_by_dept
50	9	design	60
51	9	engineer	188
52	9	marketing	124
53	9	sales	113

54	10	customer_service	336
55	10	data_science	109
56	10	design	41
57	10	engineer	171
58	10	marketing	96
59	10	sales	111
60	11	customer_service	6
61	11	data_science	2
62	11	engineer	6
63	11	marketing	2
64	12	customer_service	12
65	12	data_science	4
66	12	design	1
67	12	engineer	4
68	12	marketing	1
69	12	sales	2

In []: ##### A few quick visuals to look at how some of the features are potentially rel

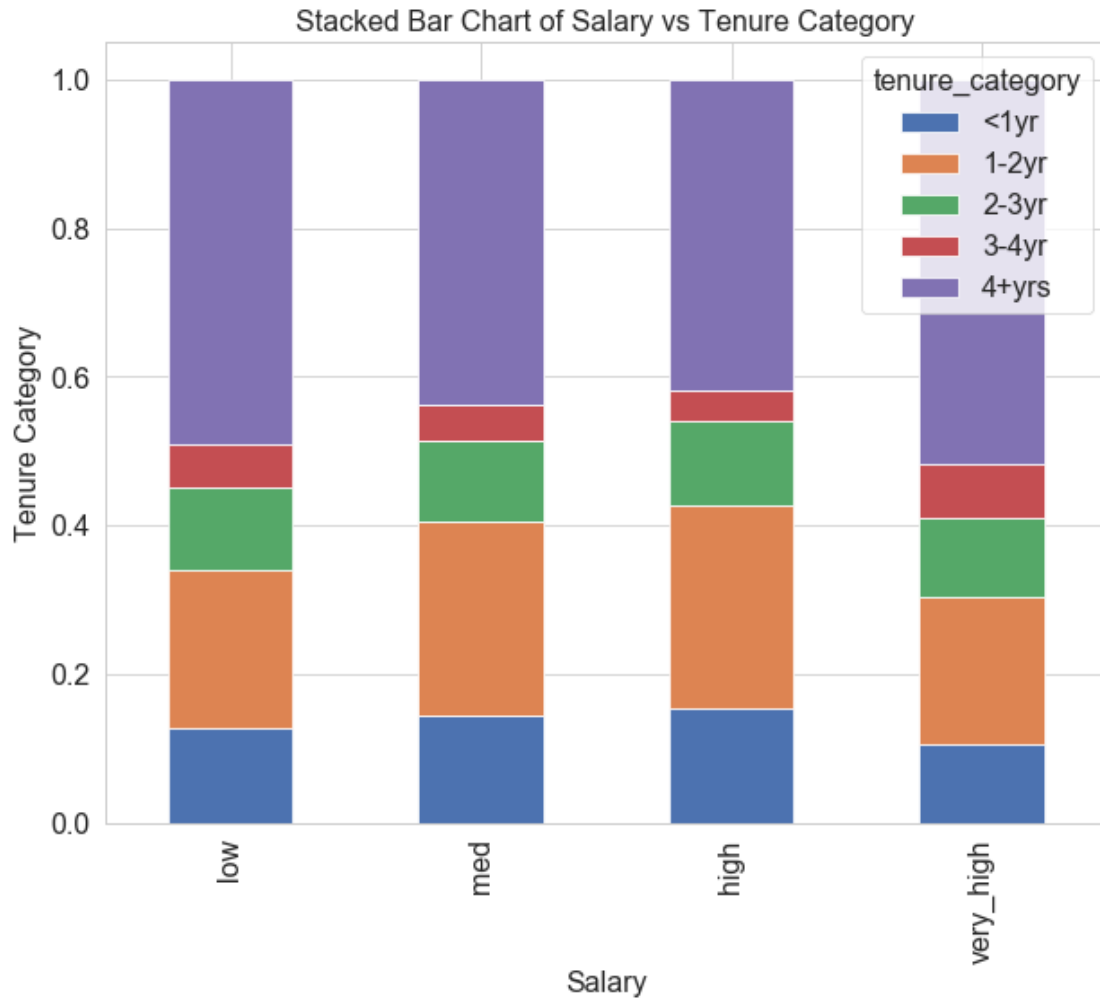
In [45]: # Salary and tenure category:

```
table=pd.crosstab(dc1_NEW.salary_rank, dc1_NEW.tenure_category)
table.div(table.sum(1).astype(float), axis=0).plot(kind='bar', stacked=True)
plt.title('Stacked Bar Chart of Salary vs Tenure Category')
plt.xlabel('Salary')
plt.ylabel('Tenure Category')
```

CONCLUDE:

Really, not stand-out trends. But, it is surprising that the very high salary group


Out[45]: Text(0, 0.5, 'Tenure Category')



```
In [46]: # Seniority and tenure category:
table=pd.crosstab(dc1_NEW.seniority, dc1_NEW.tenure_category)
table.div(table.sum(1).astype(float), axis=0).plot(kind='bar', stacked=True)
plt.title('Stacked Bar Chart of Seniority vs Tenure Category')
plt.xlabel('Seniority')
plt.ylabel('Tenure Category')

# CONCLUDE:
# No major stand-out trends.
```

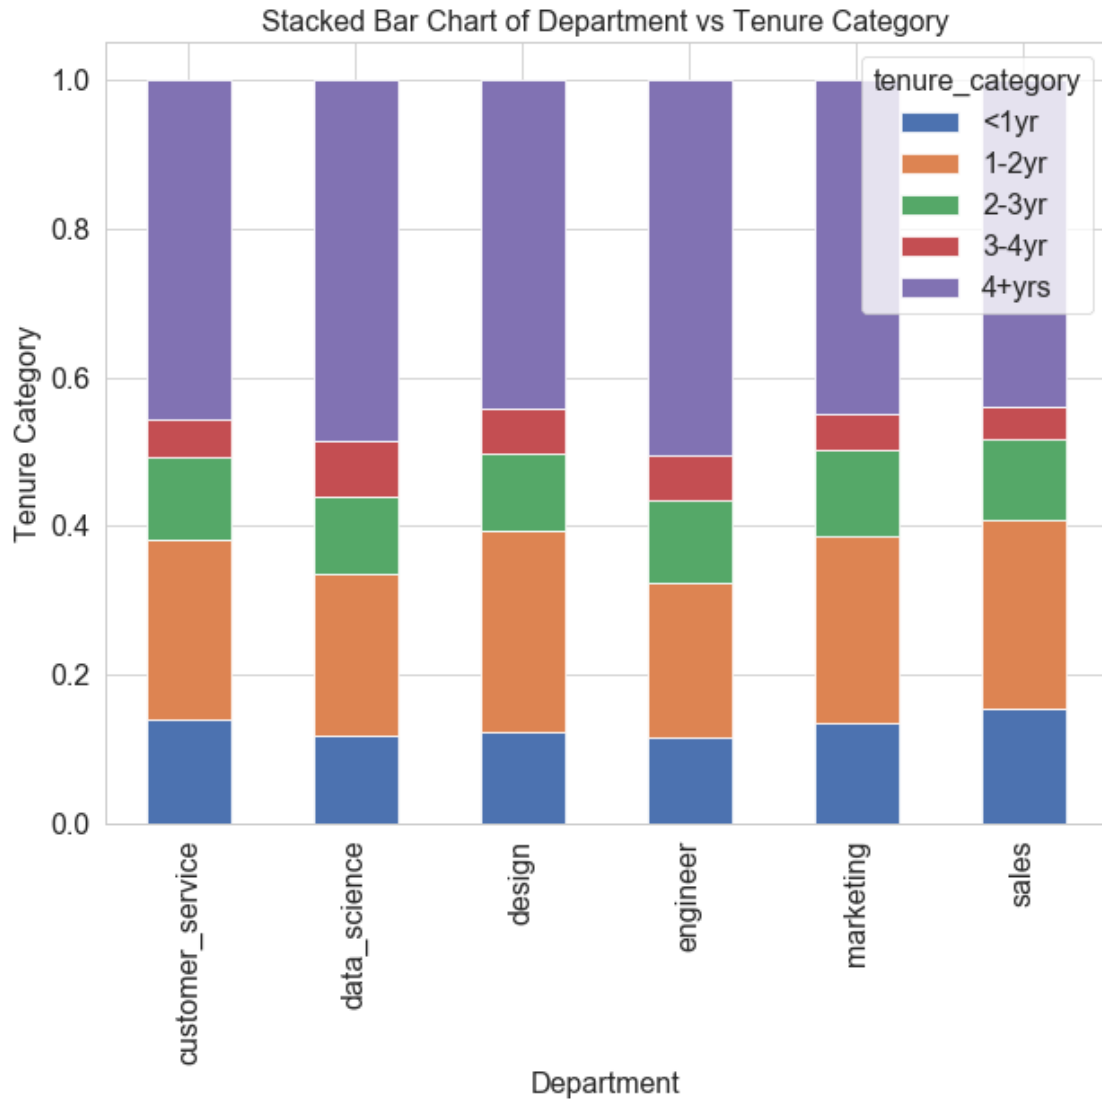
```
Out[46]: Text(0, 0.5, 'Tenure Category')
```



```
In [47]: # Department and tenure category:
table=pd.crosstab(dc1_NEW.dept, dc1_NEW.tenure_category)
table.div(table.sum(1).astype(float), axis=0).plot(kind='bar', stacked=True)
plt.title('Stacked Bar Chart of Department vs Tenure Category')
plt.xlabel('Department')
plt.ylabel('Tenure Category')

# CONCLUDE:
# No real trends here either.
```


```
Out[47]: Text(0, 0.5, 'Tenure Category')
```



```
In [48]: # Company ID and tenure category:
table=pd.crosstab(dc1_NEW.company_id, dc1_NEW.tenure_category)
table.div(table.sum(1).astype(float), axis=0).plot(kind='bar', stacked=True)
plt.title('Stacked Bar Chart of Company ID vs Tenure Category')
plt.xlabel('Company')
plt.ylabel('Tenure Category')

# CONCLUDE:
# company 11 and 12 really appear different than the others.
# previously looked at sum of number of employees - counts - need to include this her
```

```
Out[48]: Text(0, 0.5, 'Tenure Category')
```



```
In [51]: # Just to get some real counts for eyeball comparison of employment (y/n) across departments
counts = dc1_NEW.groupby(['still_employed'])['dept'].value_counts().values
counts
# what is returned is whether someone is still employed : 0 or 1 in that order for each department
# customer service: 5094 not employed; 2355 employed... and so on.
```

```
Out[51]: array([5094, 2355, 1811, 1782, 1681, 778, 4086, 2245, 1499, 1384, 1361, 602])
```

```
In [238]: # Now, let's quickly take a look JUST at the people who quit:
```

```
past_employees_df = dc1_NEW[dc1_NEW.still_employed == '0']
#past_employees_df.head(20)

# also check the opposite
#current_employees_df = dc1_NEW[dc1_NEW.still_employed == '1']
#current_employees_df.head(20)
```

```
In [62]: # Change data types for next step visualizations
past_employees_df["company_id"] = past_employees_df["company_id"].astype(int)
```

```
past_employees_df["days_employeed"] = past_employees_df["days_employeed"].astype(int)
past_employees_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 13501 entries, 0 to 24699
Data columns (total 15 columns):
employee_id      13501 non-null float64
company_id       13501 non-null int64
dept             13501 non-null object
seniority         13501 non-null int64
salary           13501 non-null float64
join_date        13501 non-null datetime64[ns]
quit_date        13501 non-null datetime64[ns]
total_tenure     13501 non-null timedelta64[ns]
days_employeed  13501 non-null int64
start_year       13501 non-null int64
start_month      13501 non-null int64
days_employeed_cont 13501 non-null float64
tenure_category  13501 non-null category
still_employeed  13501 non-null object
salary_rank      13501 non-null category
dtypes: category(2), datetime64[ns](2), float64(3), int64(5), object(2), timedelta64[ns](1)
memory usage: 2.1+ MB
```

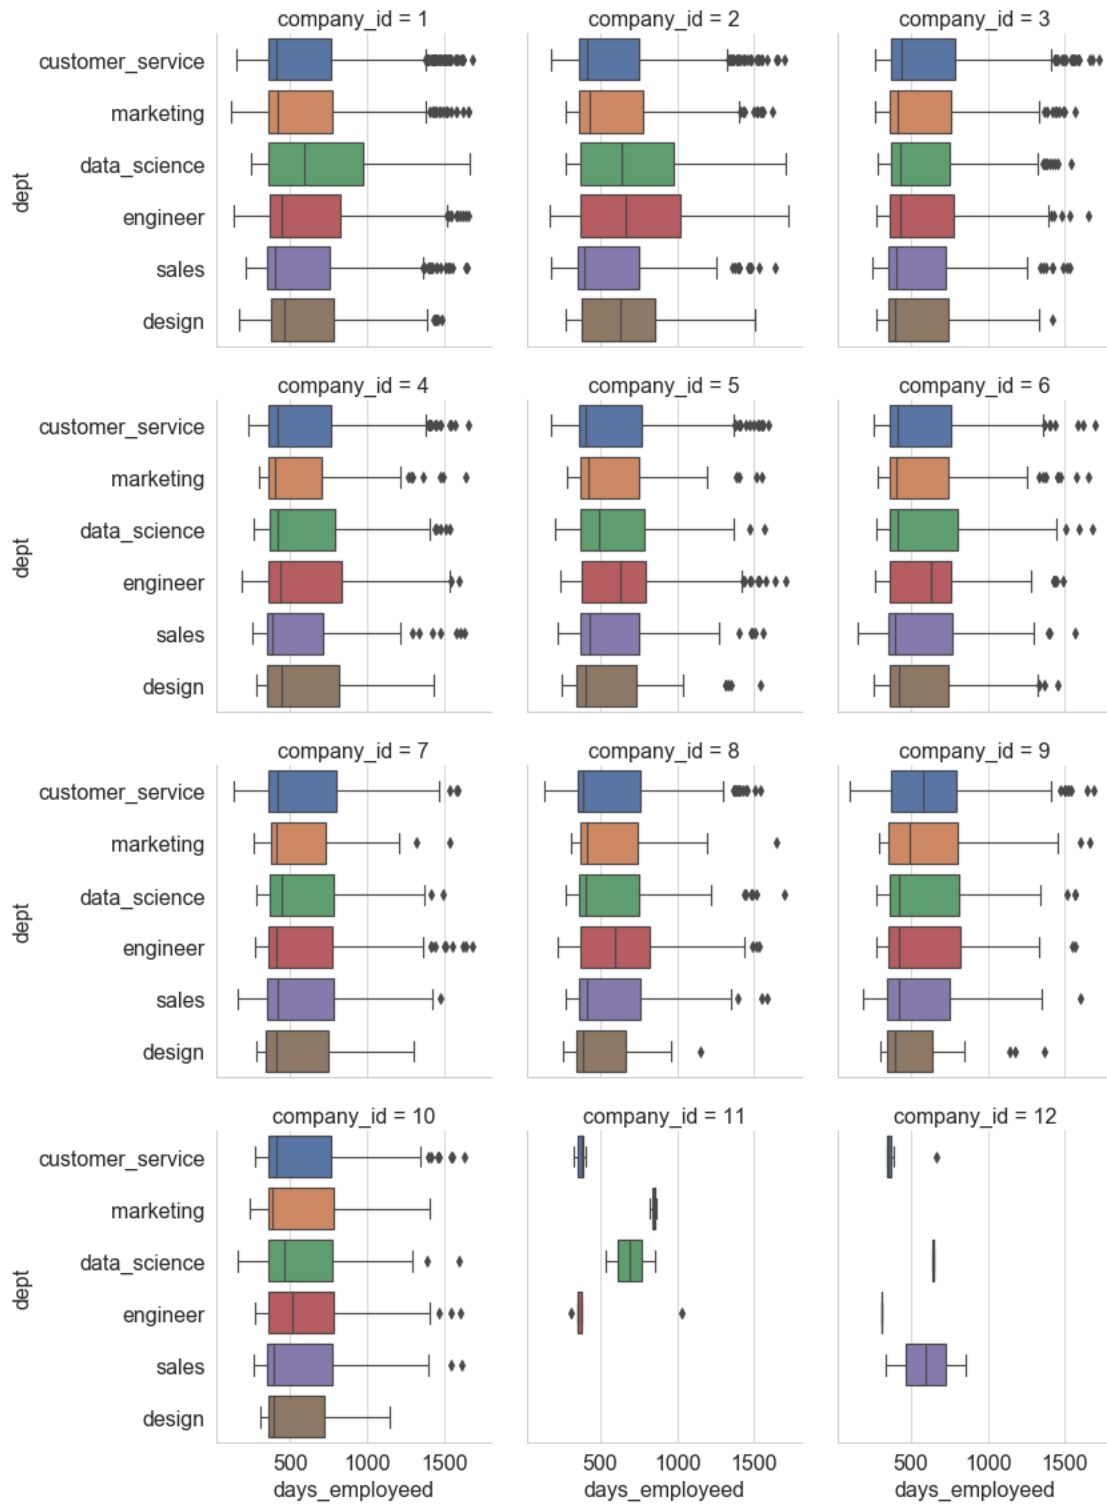
```
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html
    """Entry point for launching an IPython kernel.
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html
```

```
In [63]: # Let's visualize days employeed by department and company for this past employees da
g = sns.catplot(x="days_employeed", y="dept", col="company_id", data=past_employees_d

# CONCLUDE:
# For the people who *have* quit, it doesn't look like there's a huge departmental d
# Company 11 and 12 are really different than the rest b/c of the dept that they have
```

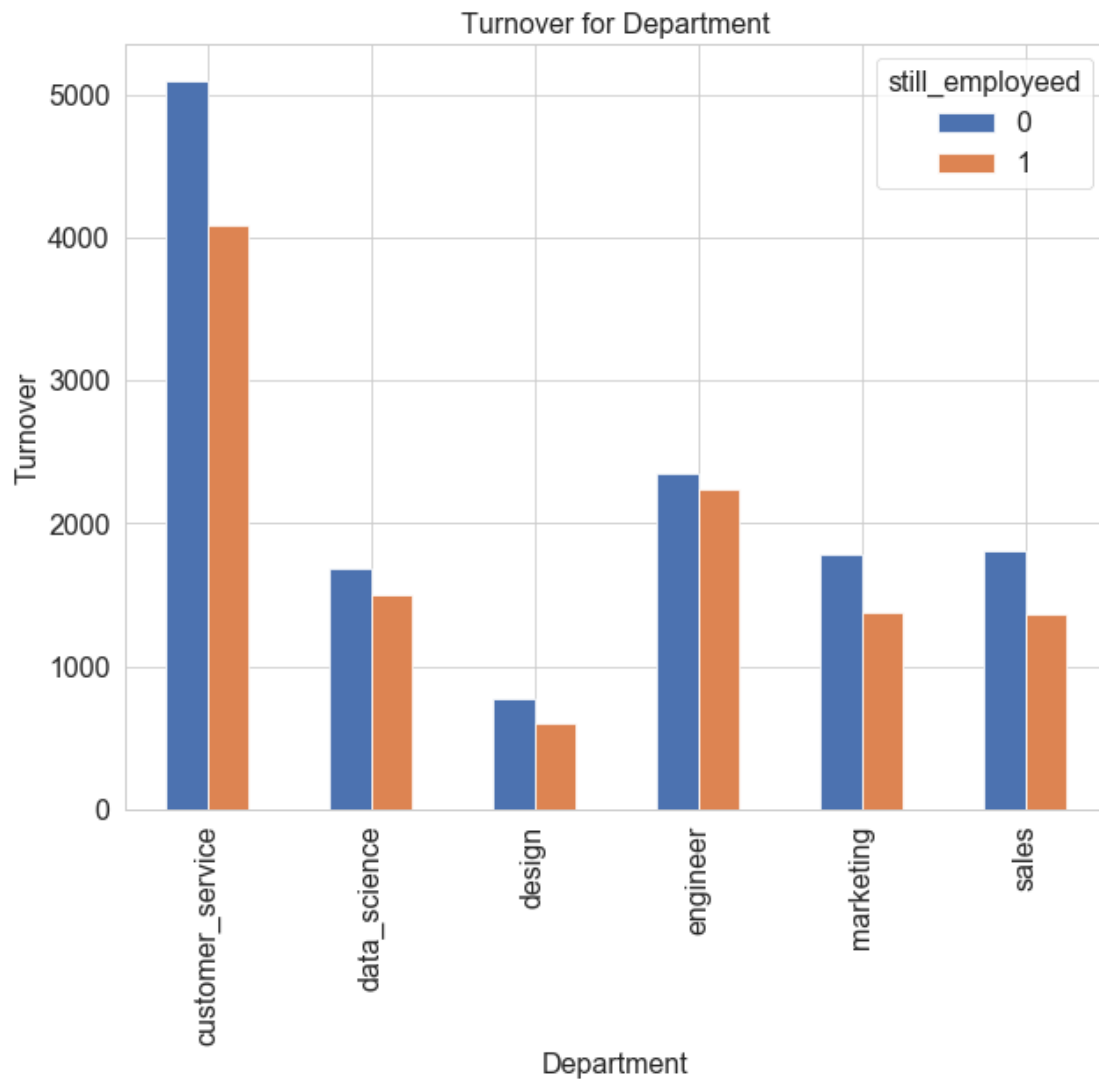
In [49]: *# Switching back to large/whole dataset, visualize relationship between dept and binary variable 'still_employed'*
 pd.crosstab(dc1_NEW.dept, dc1_NEW.still_employed).plot(kind='bar')

```
plt.title('Turnover for Department')
plt.xlabel('Department')
plt.ylabel('Turnover')
```

CONCLUDE:

Maybe, the gap between those who have left (0) and those who have stayed (1) at the

```
Out[49]: Text(0, 0.5, 'Turnover')
```

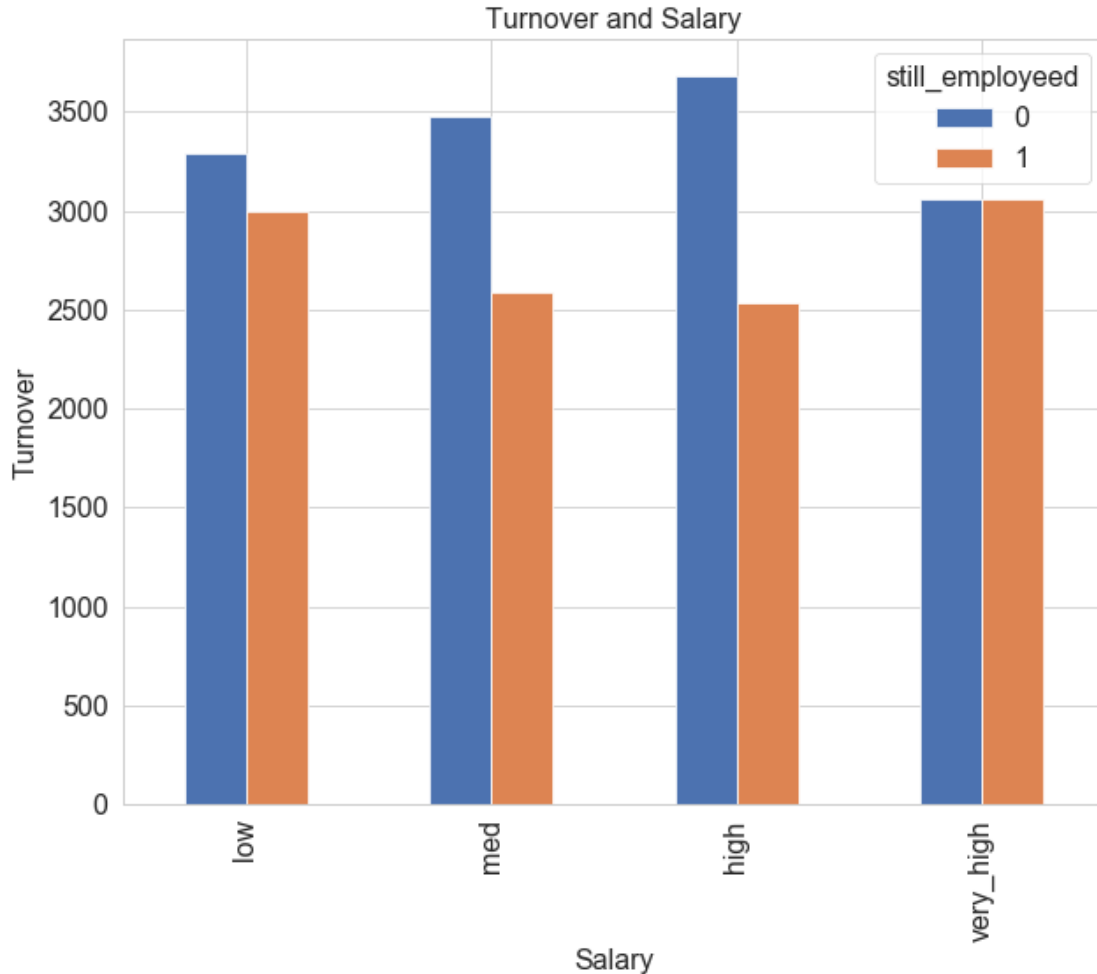


```
In [50]: # Salary and turnover classification (still employed or not)
pd.crosstab(dc1_NEW.salary_rank,dc1_NEW.still_employed).plot(kind='bar')
plt.title('Turnover and Salary')
plt.xlabel('Salary')
plt.ylabel('Turnover')
```

```
# CONCLUDE:
```

```
# Maybe, the gap between those who quit and those who are still working is larger for
```

```
Out[50]: Text(0, 0.5, 'Turnover')
```



```
In [235]: # Seniority and turnover classification (still employed or not)
```

```
plt.rcParams['figure.figsize'] = (12.0, 12.0)
```

```
pd.crosstab(dc1_NEW.seniority, dc1_NEW.still_employed).plot(kind='bar')
```

```
plt.title('Turnover and Seniority')
```

```
plt.xlabel('Seniority')
```


```
plt.ylabel('Turnover')
```

```
# CONCLUDE:
```

```
# Not a super striking pattern, but maybe for mid-range seniority levels 5-20 years,
```

```
# Only instance where more people stayed than left is 4 years experience.
```

```
Out[235]: Text(0, 0.5, 'Turnover')
```



In [65]: *# OK! Now we've looked at the data and are ready to get started with analyses with ei*
A categorical Y/outcome could be the multiple tenure categories or the binary tenur
dc1_NEW.head(20)

Out [65]:

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	
5	509529.0	4	data_science	14	165000.0	2012-01-30	
6	88600.0	4	customer_service	21	107000.0	2013-10-21	
7	716309.0	2	customer_service	4	30000.0	2014-03-05	

8	172999.0	9	engineer	7	160000.0	2012-12-10
9	504159.0	1	sales	7	104000.0	2012-06-12
10	892155.0	6	customer_service	13	72000.0	2012-11-12
11	904158.0	2	marketing	17	230000.0	2015-05-11
12	939058.0	1	marketing	1	48000.0	2012-12-10
13	163427.0	10	marketing	23	154000.0	2012-06-18
14	461248.0	2	sales	20	201000.0	2013-09-16
15	265226.0	1	data_science	4	80000.0	2014-05-27
16	932790.0	7	marketing	10	88000.0	2011-11-30
17	69693.0	7	customer_service	6	54000.0	2014-03-31
18	721600.0	2	marketing	20	193000.0	2014-12-29
19	982668.0	1	customer_service	14	76000.0	2015-07-27

	quit_date	total_tenure	days_employeed	start_year	start_month	\
0	2015-10-30	585 days	585	2014		3
1	2014-04-04	340 days	340	2013		4
2	NaT	NaT	currently_employeed	2014		10
3	2013-06-07	389 days	389	2012		5
4	2014-08-22	1040 days	1040	2011		10
5	2013-08-30	578 days	578	2012		1
6	NaT	NaT	currently_employeed	2013		10
7	NaT	NaT	currently_employeed	2014		3
8	2015-10-23	1047 days	1047	2012		12
9	NaT	NaT	currently_employeed	2012		6
10	2015-02-27	837 days	837	2012		11
11	NaT	NaT	currently_employeed	2015		5
12	2013-11-15	340 days	340	2012		12
13	2015-09-25	1194 days	1194	2012		6
14	2014-08-22	340 days	340	2013		9
15	2015-07-10	409 days	409	2014		5
16	2013-11-22	723 days	723	2011		11
17	NaT	NaT	currently_employeed	2014		3
18	NaT	NaT	currently_employeed	2014		12
19	NaT	NaT	currently_employeed	2015		7

	days_employeed_cont	tenure_category	still_employeed	salary_rank
0	585.0	1-2yr	0	med
1	340.0	<1yr	0	high
2	1850.0	4+yrs	1	med
3	389.0	1-2yr	0	med
4	1040.0	2-3yr	0	very_high
5	578.0	1-2yr	0	high
6	1850.0	4+yrs	1	med
7	1850.0	4+yrs	1	low
8	1047.0	2-3yr	0	high
9	1850.0	4+yrs	1	med
10	837.0	2-3yr	0	low
11	1850.0	4+yrs	1	very_high

12	340.0	<1yr	0	low
13	1194.0	3-4yr	0	high
14	340.0	<1yr	0	very_high
15	409.0	1-2yr	0	med
16	723.0	2-3yr	0	med
17	1850.0	4+yrs	1	low
18	1850.0	4+yrs	1	very_high
19	1850.0	4+yrs	1	low

In [73]: # I have a dataframe called 'dc1_NEW' that includes all 12 companies.

But, because company 11 and 12 seem to be very different from the rest, I may want

```
reduced_companies_df = dc1_NEW[dc1_NEW.company_id <= 10]
```

```
reduced_companies_df.head(20)
```

```
#reduced_companies_df['company_id'].unique() # double check that this did what I asked
```

```
Out[73]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	
5	509529.0	4	data_science	14	165000.0	2012-01-30	
6	88600.0	4	customer_service	21	107000.0	2013-10-21	
7	716309.0	2	customer_service	4	30000.0	2014-03-05	
8	172999.0	9	engineer	7	160000.0	2012-12-10	
9	504159.0	1	sales	7	104000.0	2012-06-12	
10	892155.0	6	customer_service	13	72000.0	2012-11-12	
11	904158.0	2	marketing	17	230000.0	2015-05-11	
12	939058.0	1	marketing	1	48000.0	2012-12-10	
13	163427.0	10	marketing	23	154000.0	2012-06-18	
14	461248.0	2	sales	20	201000.0	2013-09-16	
15	265226.0	1	data_science	4	80000.0	2014-05-27	
16	932790.0	7	marketing	10	88000.0	2011-11-30	
17	69693.0	7	customer_service	6	54000.0	2014-03-31	
18	721600.0	2	marketing	20	193000.0	2014-12-29	
19	982668.0	1	customer_service	14	76000.0	2015-07-27	

	quit_date	total_tenure	days_employeed	start_year	start_month	\
0	2015-10-30	585 days	585	2014	3	
1	2014-04-04	340 days	340	2013	4	
2	NaT	NaT	currently_employeed	2014	10	
3	2013-06-07	389 days	389	2012	5	
4	2014-08-22	1040 days	1040	2011	10	
5	2013-08-30	578 days	578	2012	1	
6	NaT	NaT	currently_employeed	2013	10	
7	NaT	NaT	currently_employeed	2014	3	

8	2015-10-23	1047 days	1047	2012	12
9	NaT	NaT	currently_employed	2012	6
10	2015-02-27	837 days	837	2012	11
11	NaT	NaT	currently_employed	2015	5
12	2013-11-15	340 days	340	2012	12
13	2015-09-25	1194 days	1194	2012	6
14	2014-08-22	340 days	340	2013	9
15	2015-07-10	409 days	409	2014	5
16	2013-11-22	723 days	723	2011	11
17	NaT	NaT	currently_employed	2014	3
18	NaT	NaT	currently_employed	2014	12
19	NaT	NaT	currently_employed	2015	7

	days_employed_cont	tenure_category	still_employed	salary_rank
0	585.0	1-2yr	0	med
1	340.0	<1yr	0	high
2	1850.0	4+yrs	1	med
3	389.0	1-2yr	0	med
4	1040.0	2-3yr	0	very_high
5	578.0	1-2yr	0	high
6	1850.0	4+yrs	1	med
7	1850.0	4+yrs	1	low
8	1047.0	2-3yr	0	high
9	1850.0	4+yrs	1	med
10	837.0	2-3yr	0	low
11	1850.0	4+yrs	1	very_high
12	340.0	<1yr	0	low
13	1194.0	3-4yr	0	high
14	340.0	<1yr	0	very_high
15	409.0	1-2yr	0	med
16	723.0	2-3yr	0	med
17	1850.0	4+yrs	1	low
18	1850.0	4+yrs	1	very_high
19	1850.0	4+yrs	1	low

In []: *# Now I also have a dataframe called 'reduced_companies_df' that includes just companies*

In []: *##### MODELING PORTION OF NB STARTS HERE ##### --->*

In [97]: *# I will first define the X array and the Y/target variable in two sets from the reduced_companies_df*
'Can I determine whether someone stays at a company for less than a year, 1-2 years or more'?

```
y = reduced_companies_df[['tenure_category']]
```

I decided not to include company ID as a predictor here since I presumably want to use company ID as a target variable
will use continuous values of salary here, first instead of the categorical salary_rank

```
X = reduced_companies_df[['dept', 'seniority', 'salary', 'start_year', 'start_month']]
```

```
In [98]: # Check new X df
X.head()
#X.shape
```

```
Out[98]:
```

	dept	seniority	salary	start_year	start_month
0	customer_service	28	89000.0	2014	3
1	marketing	20	183000.0	2013	4
2	marketing	14	101000.0	2014	10
3	customer_service	20	115000.0	2012	5
4	data_science	23	276000.0	2011	10

```
In [99]: # Check new y dataframe
y.head()
y.shape
```

```
Out[99]: (24638, 1)
```

```
In [100]: # Get dummy variables for categorical features; use drop_first to get rid of the ext
X = pd.concat([X,pd.get_dummies(X['dept'], prefix='dept', drop_first=True)],axis=1)
X.head()

# drop old dept column since it is no longer informative.
X = X.drop(['dept'], axis=1)
X.head()
```

```
Out[100]:
```

	seniority	salary	start_year	start_month	dept_data_science \
0	28	89000.0	2014	3	0
1	20	183000.0	2013	4	0
2	14	101000.0	2014	10	0
3	20	115000.0	2012	5	0
4	23	276000.0	2011	10	1

	dept_design	dept_engineer	dept_marketing	dept_sales
0	0	0	0	0
1	0	0	1	0
2	0	0	1	0
3	0	0	0	0
4	0	0	0	0

```
In [101]: #check new X dataframe
X.shape
```

```
Out[101]: (24638, 9)
```

```
In [102]: # Import modules and libraries for modeling

from sklearn.model_selection import train_test_split # to create the split in training
from sklearn.model_selection import RandomizedSearchCV # test a bunch of parameter values
from sklearn.model_selection import GridSearchCV # hone in on exact parameter values
```



```

from sklearn.model_selection import cross_val_score # to get performance in the 5 (o
from sklearn import metrics # to be able to use recall and any other metric for mode
from sklearn.metrics import classification_report # nice format to show all model pe
from sklearn.metrics import confusion_matrix # classic TP/FP assessment
from sklearn.metrics import roc_curve # another metric of performance for comparing

In [103]: # create training and test data sets based on X and y defined above. Will set test s

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_st

In [104]: # take a look at the new X_train, X_test, y_train, and y_test
X_train.head()
X_train.shape

Out[104]: (19710, 9)

In [105]: X_test.shape

Out[105]: (4928, 9)

In [106]: y_train.shape

Out[106]: (19710, 1)

In [107]: y_test.shape
y_test.head()

Out[107]:
      tenure_category
4396                <1yr
20473               2-3yr
9680                4+yrs
21359               4+yrs
15854               4+yrs

In [108]: # Check for class imbalance, which could affect model performance metrics like accur
y_train.tenure_category.value_counts()

Out[108]: 4+yrs      9181
          1-2yr     4654
          <1yr      2625
          2-3yr     2170
          3-4yr     1080
          Name: tenure_category, dtype: int64

In [109]: y_test.tenure_category.value_counts()

Out[109]: 4+yrs      2296
          1-2yr     1163
          <1yr       656
          2-3yr      543
          3-4yr      270
          Name: tenure_category, dtype: int64

```

```
In [110]: # Determine that there is a pretty big class imbalance: i will use smote to balance
```

```
import imblearn as imblearn # for SMOTE applied to training set only
from imblearn.over_sampling import SMOTE # for SMOTE applied to training set only
from collections import Counter # to easily check old and new class sizes
```

```
sm = SMOTE(random_state=11)
X_res_train, y_res_train = sm.fit_resample(X_train, y_train)
print('Resampled dataset shape %s' % Counter(y_res_train))
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/utils/validation.py:761: DataConversionWarning:
  y = column_or_1d(y, warn=True)
```

```
Resampled dataset shape Counter({'1-2yr': 9181, '4+yrs': 9181, '<1yr': 9181, '2-3yr': 9181, '3+yrs': 9181})
```

```
In [114]: # Check new df
          y_res_train.shape
          #y_res_train
```

```
Out[114]: (45905,)
```

```
In [116]: # Check new df
          X_res_train.shape
          #X_res_train
```

```
Out[116]: (45905, 9)
```

```
In [ ]: # Great! Now I'm ready to run some classifier models. Can't use logistic regression wi
        # I won't do any recursive feature selection to begin because there's already such a r
        # Start by defining some randomized search parameters for hyperparameter tuning follow
```

```
In [144]: # parameters that CAN be set in RF classifier; set up random space over which to sea
```

```
# Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 10, stop = 100, num = 5)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt', 2, 4, 8]
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(2, 10, 2)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [10, 50]
# Minimum number of samples required at each leaf node
min_samples_leaf = [5, 10, 25]
# Method of selecting samples for training each tree
#bootstrap = [True]
```

```

In [145]: # Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}

print(random_grid)

{'n_estimators': [10, 32, 55, 77, 100], 'max_features': ['auto', 'sqrt', 2, 4, 8], 'max_depth':

In [146]: from sklearn.ensemble import RandomForestClassifier

# Use the random grid to search for best hyperparameters
# First create the base model to tune
RF = RandomForestClassifier()
#RF.fit(X_res_train, y_res_train)
#y_pred=RF.predict(X_test)
#print(classification_report(y_test, y_pred))

# Random search of parameters, using 5 fold cross validation, on the training set; I
RF_random = RandomizedSearchCV(estimator = RF, param_distributions = random_grid, n_

In [147]: # Fit the random search model to the SMOTE adjusted training data set
RF_random.fit(X_res_train, y_res_train)

Fitting 5 folds for each of 100 candidates, totalling 500 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 42.3s
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 3.8min
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 7.2min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 9.8min finished

Out[147]: RandomizedSearchCV(cv=5, error_score='raise-deprecating',
                             estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criter
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators='warn', n_jobs=None,
                             oob_score=False, random_state=None, verbose=0,
                             warm_start=False),
                             fit_params=None, iid='warn', n_iter=100, n_jobs=-1,
                             param_distributions={'n_estimators': [10, 32, 55, 77, 100], 'max_features'

```

```
pre_dispatch='2*n_jobs', random_state=42, refit=True,  
return_train_score='warn', scoring='accuracy', verbose=2)
```

```
In [148]: # Get the best parameters from the randomized search  
RF_random.best_params_
```

```
Out[148]: {'n_estimators': 100,  
          'min_samples_split': 10,  
          'min_samples_leaf': 5,  
          'max_features': 8,  
          'max_depth': None,  
          'bootstrap': True}
```

```
In [149]: # Run cross validation using the best parameters and then average them to get an ide  
CV_scores = cross_val_score(RF_random, X_res_train, y_res_train, cv=5, scoring = 'ac  
print(CV_scores)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 40.9s  
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 2.5min  
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 5.2min  
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 7.0min finished
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 30.1s  
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 2.2min  
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 4.7min  
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 6.8min finished
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 29.0s  
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 2.1min  
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 4.5min  
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 6.1min finished
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 27.7s
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 2.0min
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 4.3min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 5.9min finished
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 27.5s
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 2.0min
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 4.3min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 5.9min finished
```

```
[0.43048449 0.60141612 0.64858388 0.61448802 0.61023965]
```

```
In [150]: print("Accuracy of A RF classifier in Cross Validation: %0.2f (+/- %0.2f)" % (CV_score, CV_std))
Accuracy of A RF classifier in Cross Validation: 0.58 (+/- 0.15)
```

```
In [151]: print(classification_report(y_test, y_pred))
           # Not very good performance but notice that the model has an easier time classifying
           # Maybe it would be easier to lump everyone who has quit (groups <1year up to 3-4years)
```

	precision	recall	f1-score	support
1-2yr	0.33	0.37	0.35	1163
2-3yr	0.20	0.20	0.20	543
3-4yr	0.20	0.24	0.22	270
4+yrs	0.71	0.70	0.70	2296
<1yr	0.18	0.12	0.14	656
micro avg	0.46	0.46	0.46	4928
macro avg	0.32	0.33	0.32	4928
weighted avg	0.46	0.46	0.46	4928

```
In [152]: # Showign the breakdown of classification with this RF model.
           confusion_matrix(y_test, y_pred)
```

```
Out[152]: array([[ 432,  168,  106,  315,  142],
                 [ 222,  111,   60,   97,   53],
                 [  96,   62,   65,   21,   26],
                 [ 364,  125,   53, 1604,  150],
                 [ 207,   89,   48,  233,   79]])
```

```
In [153]: print('Accuracy of A RF classifier on Test Set:', metrics.accuracy_score(y_test,y_pr
```

Accuracy of A RF classifier on Test Set: 0.4648944805194805

```
In [162]: # Re-run the model with parameters that I found above in randomized search:
```

```
# {'n_estimators': 100,  
# 'min_samples_split': 10,  
# 'min_samples_leaf': 5,  
# 'max_features': 8,  
# 'max_depth': None,  
# 'bootstrap': True}
```

```
RF_test = RandomForestClassifier(n_estimators=100, min_samples_split=10, min_samples_  
#RF_test = RandomForestClassifier()  
RF_test.fit(X_res_train, y_res_train)  
y_pred=RF_test.predict(X_test)  
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
1-2yr	0.33	0.37	0.35	1163
2-3yr	0.20	0.20	0.20	543
3-4yr	0.21	0.32	0.25	270
4+yrs	0.72	0.74	0.73	2296
<1yr	0.18	0.08	0.11	656
micro avg	0.48	0.48	0.48	4928
macro avg	0.33	0.34	0.33	4928
weighted avg	0.47	0.48	0.47	4928

```
In [163]: # For this poorly performing model we can get an idea of which features might matter  
names = X.columns  
feature_imp = pd.Series(RF_test.feature_importances_,index=names).sort_values(ascending=  
feature_imp
```

```
Out[163]: start_year      0.390510  
salary      0.224351  
seniority    0.146317  
start_month  0.141169  
dept_sales   0.023781  
dept_marketing 0.021610  
dept_engineer 0.019929  
dept_data_science 0.019435  
dept_design  0.012898  
dtype: float64
```

```
In [164]: # Plot feature importance values from the RF classifier.
```

```
from matplotlib import rcParams
```

```
names = X.columns
```

```
feature_imp = pd.Series(RF_test.feature_importances_, index=names).sort_values(ascending=False)  
feature_imp
```

```
plt.rcParams['figure.figsize'] = (15, 12)
```

```
plt.rcParams['axes.labelweight'] = 'bold'
```

```
#plt.rcParams["axes.labelsize"] = 30
```

```
rcParams['axes.titlepad'] = 20
```

```
sns.set(style='whitegrid')
```

```
sns.set_context("paper", rc={"font.size":32,"axes.titlesize":32,"axes.labelsize":30,
```

```
sns.barplot(x=feature_imp, y=feature_imp.index, palette='Reds_d')
```

```
# Add labels to your graph
```

```
plt.xlabel('Importance Score')
```


```
#plt.ylabel('Features')
```

```
plt.title("Visualizing Important Features from a Random Forest Classifier")
```

```
plt.legend()
```

No handles with labels found to put in legend.

```
Out[164]: <matplotlib.legend.Legend at 0x1a1dcc04e0>
```



```
In [ ]: # Summary: the RF model does not do particularly well at predicting employee tenure wi
# Maybe, it would be better to predict based on the binary - quit or stayed (e.g. are
# this means we could use an outcome variable like 'still_employed' and apply logisti
# another alternative route would be to use days employeeed as a continous outcome (e.g
```

```
In [165]: # A quick reminder about the df: I will again use the reduced set of 10 companies fo
```

```
reduced_companies_df.head(20)
```

```
Out[165]:
```

	employee_id	company_id	dept	seniority	salary	join_date	\
0	13021.0	7	customer_service	28	89000.0	2014-03-24	
1	825355.0	7	marketing	20	183000.0	2013-04-29	
2	927315.0	4	marketing	14	101000.0	2014-10-13	
3	662910.0	7	customer_service	20	115000.0	2012-05-14	
4	256971.0	2	data_science	23	276000.0	2011-10-17	
5	509529.0	4	data_science	14	165000.0	2012-01-30	
6	88600.0	4	customer_service	21	107000.0	2013-10-21	
7	716309.0	2	customer_service	4	30000.0	2014-03-05	
8	172999.0	9	engineer	7	160000.0	2012-12-10	
9	504159.0	1	sales	7	104000.0	2012-06-12	
10	892155.0	6	customer_service	13	72000.0	2012-11-12	
11	904158.0	2	marketing	17	230000.0	2015-05-11	
12	939058.0	1	marketing	1	48000.0	2012-12-10	
13	163427.0	10	marketing	23	154000.0	2012-06-18	
14	461248.0	2	sales	20	201000.0	2013-09-16	
15	265226.0	1	data_science	4	80000.0	2014-05-27	
16	932790.0	7	marketing	10	88000.0	2011-11-30	
17	69693.0	7	customer_service	6	54000.0	2014-03-31	
18	721600.0	2	marketing	20	193000.0	2014-12-29	
19	982668.0	1	customer_service	14	76000.0	2015-07-27	

	quit_date	total_tenure	days_employeeed	start_year	start_month	\
0	2015-10-30	585 days	585	2014	3	
1	2014-04-04	340 days	340	2013	4	
2	NaT	NaT	currently_employeeed	2014	10	
3	2013-06-07	389 days	389	2012	5	
4	2014-08-22	1040 days	1040	2011	10	
5	2013-08-30	578 days	578	2012	1	
6	NaT	NaT	currently_employeeed	2013	10	
7	NaT	NaT	currently_employeeed	2014	3	
8	2015-10-23	1047 days	1047	2012	12	
9	NaT	NaT	currently_employeeed	2012	6	
10	2015-02-27	837 days	837	2012	11	
11	NaT	NaT	currently_employeeed	2015	5	
12	2013-11-15	340 days	340	2012	12	

13	2015-09-25	1194 days	1194	2012	6
14	2014-08-22	340 days	340	2013	9
15	2015-07-10	409 days	409	2014	5
16	2013-11-22	723 days	723	2011	11
17	NaT	NaT	currently_employeed	2014	3
18	NaT	NaT	currently_employeed	2014	12
19	NaT	NaT	currently_employeed	2015	7

	days_employeed_cont	tenure_category	still_employeed	salary_rank
0	585.0	1-2yr	0	med
1	340.0	<1yr	0	high
2	1850.0	4+yrs	1	med
3	389.0	1-2yr	0	med
4	1040.0	2-3yr	0	very_high
5	578.0	1-2yr	0	high
6	1850.0	4+yrs	1	med
7	1850.0	4+yrs	1	low
8	1047.0	2-3yr	0	high
9	1850.0	4+yrs	1	med
10	837.0	2-3yr	0	low
11	1850.0	4+yrs	1	very_high
12	340.0	<1yr	0	low
13	1194.0	3-4yr	0	high
14	340.0	<1yr	0	very_high
15	409.0	1-2yr	0	med
16	723.0	2-3yr	0	med
17	1850.0	4+yrs	1	low
18	1850.0	4+yrs	1	very_high
19	1850.0	4+yrs	1	low

```
In [166]: # Here the outcome/predicted variable will be binary, still employeed or not.
y = reduced_companies_df[['still_employeed']]
```

```
# I decided not to include company ID as a predictor here since I presumably want to
# will use continuous values of salary here, first instead of the categorical salary
```

```
X = reduced_companies_df[['dept', 'seniority', 'salary', 'start_year', 'start_month']]
```

```
In [167]: y.head()
```

```
Out[167]: still_employeed
0          0
1          0
2          1
3          0
4          0
```

```
In [168]: y.shape
```

```
Out[168]: (24638, 1)
```

```
In [169]: X.head()
```

```
Out[169]:
```

	dept	seniority	salary	start_year	start_month
0	customer_service	28	89000.0	2014	3
1	marketing	20	183000.0	2013	4
2	marketing	14	101000.0	2014	10
3	customer_service	20	115000.0	2012	5
4	data_science	23	276000.0	2011	10

```
In [170]: # get dummy variables; use drop_first to get rid of the extra column when creating dummies
X = pd.concat([X, pd.get_dummies(X['dept'], prefix='dept', drop_first=True)], axis=1)
X.head()
```

```
# drop old dept column since it is no longer informative.
X = X.drop(['dept'], axis=1)
X.head()
```

```
Out[170]:
```

	seniority	salary	start_year	start_month	dept_data_science \
0	28	89000.0	2014	3	0
1	20	183000.0	2013	4	0
2	14	101000.0	2014	10	0
3	20	115000.0	2012	5	0
4	23	276000.0	2011	10	1

	dept_design	dept_engineer	dept_marketing	dept_sales
0	0	0	0	0
1	0	0	1	0
2	0	0	1	0
3	0	0	0	0
4	0	0	0	0

```
In [171]: # create training and test data sets based on X and y defined above. Will set test size to 0.2
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=11)
```

```
In [173]: # check for class imbalance
y_train.still_employed.value_counts()
```

```
Out[173]: 0    10781
          1     8929
          Name: still_employed, dtype: int64
```

```
In [175]: # Class imbalance is not too bad here, but I will apply SMOTE anyway, to give us the same number of samples for both classes
sm = SMOTE(random_state=11)
X_res_train, y_res_train = sm.fit_resample(X_train, y_train.values.ravel())
print('Resampled dataset shape %s' % Counter(y_res_train))
```

```
Resampled dataset shape Counter({'1': 10781, '0': 10781})
```

```
In [176]: # Now we are ready now to try logistic regression and then RF with this binary outcome
```

```
RF = RandomForestClassifier()
RF.fit(X_res_train, y_res_train)
y_pred=RF.predict(X_test)
print(classification_report(y_test, y_pred))
```

```
# Not too bad! no parameter tuning at this point and already doing better than the p
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/ensemble/forest.py:246: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

	precision	recall	f1-score	support
0	0.75	0.81	0.78	2696
1	0.75	0.67	0.71	2232
micro avg	0.75	0.75	0.75	4928
macro avg	0.75	0.74	0.75	4928
weighted avg	0.75	0.75	0.75	4928

```
In [177]: confusion_matrix(y_test, y_pred)
```

```
Out[177]: array([[2194,  502],
                 [ 726, 1506]])
```

```
In [179]: print('Overall accuracy of a binary RF classifier on the Test Set:', metrics.accuracy_score(y_test, y_pred))
```

```
Overall accuracy of a binary RF classifier on the Test Set: 0.7508116883116883
```

```
In [180]: names = X.columns
feature_imp = pd.Series(RF.feature_importances_, index=names).sort_values(ascending=False)
feature_imp
```

```
Out[180]: start_year      0.481252
salary      0.259028
seniority    0.141801
start_month  0.086294
dept_marketing 0.007004
dept_engineer 0.006501
dept_sales   0.006479
dept_data_science 0.006001
dept_design  0.005639
dtype: float64
```

```
In [181]: names = X.columns
```


```
feature_imp = pd.Series(RF.feature_importances_, index=names).sort_values(ascending=False)
feature_imp
```

```
plt.rcParams['figure.figsize'] = (15, 12)
plt.rcParams['axes.labelweight'] = 'bold'
#plt.rcParams["axes.labelsize"] = 30
rcParams['axes.titlepad'] = 20
sns.set(style='whitegrid')
sns.set_context("paper", rc={"font.size":32,"axes.titlesize":32,"axes.labelsize":30,

sns.barplot(x=feature_imp, y=feature_imp.index, palette='Reds_d')
# Add labels to your graph
plt.xlabel('Importance Score')
#plt.ylabel('Features')
plt.title("Visualizing Important Features from a Random Forest Classifier")
plt.legend()
```

No handles with labels found to put in legend.

```
Out[181]: <matplotlib.legend.Legend at 0x1a200bdd30>
```



```
In [195]: # Summary up to here: I could now tune the hyperparameters of this RF before moving
```

```
from sklearn.linear_model import LogisticRegression
```

```
logreg = LogisticRegression()  
#logreg.fit(X_res_train, y_res_train)  
#y_pred = logreg.predict(X_test)  
#logreg.score(X_test, y_test)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
Out[195]: 0.5355113636363636
```

```
In [186]: c_space = np.logspace(-5, 8, 15)  
param_grid = {'C': c_space, 'penalty': ['l1', 'l2']}
```

```
In [187]: # Instantiate the GridSearchCV object: logreg_cv  
logreg_cv = GridSearchCV(logreg, param_grid, cv=5, scoring = 'accuracy')
```

```
In [188]: logreg_cv.fit(X_res_train, y_res_train)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De  
FutureWarning)
```

[illegible]

[illegible]

```
/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning: Liblinear :
  "the number of iterations.", ConvergenceWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning: Liblinear :
  "the number of iterations.", ConvergenceWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning: Liblinear :
  "the number of iterations.", ConvergenceWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning: Liblinear :
  "the number of iterations.", ConvergenceWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning: Liblinear :
  "the number of iterations.", ConvergenceWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning: Liblinear :
  "the number of iterations.", ConvergenceWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
  FutureWarning)
```

```

/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning: Liblinear :
"the number of iterations.", ConvergenceWarning)

```

```

Out[188]: GridSearchCV(cv=5, error_score='raise-deprecating',
    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False),
    fit_params=None, iid='warn', n_jobs=None,
    param_grid={'C': array([1.00000e-05, 8.48343e-05, 7.19686e-04, 6.10540e-03, 5.00000e-02,
    4.39397e-01, 3.72759e+00, 3.16228e+01, 2.68270e+02, 2.27585e+03,
    1.93070e+04, 1.63789e+05, 1.38950e+06, 1.17877e+07, 1.00000e+08])}, 'penalty': 'l1',
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='accuracy', verbose=0)

```

```

In [189]: logreg_cv.best_params_

```

```

Out[189]: {'C': 268.2695795279727, 'penalty': 'l1'}

```

```

In [190]: CV_scores = cross_val_score(logreg, X_res_train, y_res_train, cv=5, scoring = 'accuracy')
print(CV_scores)

```

```

/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
FutureWarning)
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
FutureWarning)

```

```

[0.5336115  0.5403525  0.52017625 0.50834879 0.53200371]

```

```

/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: De
FutureWarning)

```

```

In [191]: print("Accuracy of A logistic regression classifier in Cross Validation: %0.2f (+/- %0.02f)" % (CV_scores.mean(), CV_scores.std()))

Accuracy of A logistic regression classifier in Cross Validation: 0.53 (+/- 0.02)

```

```


In [192]: print('Accuracy of A logistic regression classifier on the Test Set:', metrics.accuracy_score(y_test, logreg.predict(X_test)))

```

Accuracy of A logistic regression classifier on the Test Set: 0.5355113636363636

```
In [196]: coef = logreg.coef_[0]
names = X.columns
plt.rcParams['figure.figsize'] = (10.0, 8.0)
_ = plt.plot(range(len(names)), coef)
_ = plt.xticks(range(len(names)), names, rotation = 90)
_ = plt.ylabel('Coefficients From Logistic Regression Model')
plt.show()
```

start year and start month is important in both the RF and logreg models. There's a



```
In [ ]: # So it appears that the logistic regression model performs less well than the random forest.
        # I will now go back to the RF classifier and look a little bit deeper at what the features are.
```

```
In [198]: # Go back to the binary random forest classifier and do some hyperparameter tuning.
```

```
RF = RandomForestClassifier()
```

```
In [201]: # parameters that CAN be set in RF classifier:
```

```
# Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 10, stop = 1000, num = 6)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt', 2, 4, 8]
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(2, 10, 2)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [10, 50]
# Minimum number of samples required at each leaf node
min_samples_leaf = [5, 10, 25]
# Method of selecting samples for training each tree
#bootstrap = [True]
```

```
In [202]: # Create the random grid
```

```
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}
```

```
print(random_grid)
```

```
{'n_estimators': [10, 208, 406, 604, 802, 1000], 'max_features': ['auto', 'sqrt', 2, 4, 8], 'min_samples_split': [10, 50], 'min_samples_leaf': [5, 10, 25], 'bootstrap': [True]}
```

```
In [203]: # Random search of parameters, using 5 fold cross validation, on the training set
```

```
RF_random = RandomizedSearchCV(estimator = RF, param_distributions = random_grid, n_jobs = -1)
# Fit the random search model to the SMOTE adjusted training data set
RF_random.fit(X_res_train, y_res_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 33 tasks | elapsed: 2.1min
```

```
[Parallel(n_jobs=-1)]: Done 154 tasks      | elapsed: 8.1min
[Parallel(n_jobs=-1)]: Done 357 tasks      | elapsed: 16.4min
/anaconda3/lib/python3.6/site-packages/sklearn/externals/joblib/externals/loky/process_executor.py:155: UserWarning:
  "timeout or by a memory leak.", UserWarning
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 22.8min finished
```

```
Out[203]: RandomizedSearchCV(cv=5, error_score='raise-deprecating',
                             estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion='entropy',
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators='warn', n_jobs=None,
                             oob_score=False, random_state=None, verbose=0,
                             warm_start=False),
                             fit_params=None, iid='warn', n_iter=100, n_jobs=-1,
                             param_distributions={'n_estimators': [10, 208, 406, 604, 802, 1000], 'max_depth': [None, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, 120, 130, 140, 150, 160, 170, 180, 190, 200]},
                             pre_dispatch='2*n_jobs', random_state=42, refit=True,
                             return_train_score='warn', scoring='accuracy', verbose=2)
```

```
In [204]: RF_random.best_params_
```

```
Out[204]: {'n_estimators': 1000,
            'min_samples_split': 10,
            'min_samples_leaf': 5,
            'max_features': 'auto',
            'max_depth': 10,
            'bootstrap': True}
```

```
In [205]: CV_scores = cross_val_score(RF_random, X_res_train, y_res_train, cv=5, scoring = 'accuracy')
          print(CV_scores)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 1.2min
[Parallel(n_jobs=-1)]: Done 154 tasks      | elapsed: 6.5min
[Parallel(n_jobs=-1)]: Done 357 tasks      | elapsed: 12.9min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 17.7min finished
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 1.5min
[Parallel(n_jobs=-1)]: Done 154 tasks      | elapsed: 6.7min
[Parallel(n_jobs=-1)]: Done 357 tasks      | elapsed: 13.0min
```



```

/anaconda3/lib/python3.6/site-packages/sklearn/externals/joblib/externals/loky/process_executor.py:101:
  "timeout or by a memory leak.", UserWarning
/anaconda3/lib/python3.6/site-packages/sklearn/externals/joblib/externals/loky/process_executor.py:101:
  "timeout or by a memory leak.", UserWarning
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 18.0min finished

```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 1.4min
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 7.2min
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 13.0min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 17.4min finished

```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 1.8min
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 6.7min
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 13.1min
/anaconda3/lib/python3.6/site-packages/sklearn/externals/joblib/externals/loky/process_executor.py:101:
  "timeout or by a memory leak.", UserWarning
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 17.4min finished

```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 6.2min
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 11.7min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 17.5min finished

```

```

[0.78789986 0.80612245 0.78501855 0.79754174 0.81864564]

```

```

In [208]: print("Average accuracy of a tuned classifier in 5-fold cross validation on the training data set : 0.78789986")
Average accuracy of a tuned classifier in 5-fold cross validation on the training data set : 0.78789986

```

```

In [209]: print('Accuracy of A logistic regression classifier on the held out test data set:', 0.7508116883116883)
Accuracy of A logistic regression classifier on the held out test data set: 0.7508116883116883

```

```
In [ ]: # The hyperparameter tuning applied to this RF classifier on a binary outcome (quit vs
```

```
In [210]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.75	0.82	0.78	2696
1	0.75	0.67	0.71	2232
micro avg	0.75	0.75	0.75	4928
macro avg	0.75	0.74	0.75	4928
weighted avg	0.75	0.75	0.75	4928

```
In [213]: # here is a list of the best parameters for the model:
```

```
{'n_estimators': 1000,  
  'min_samples_split': 10,  
  'min_samples_leaf': 5,  
  'max_features': 'auto',  
  'max_depth': 10,  
  'bootstrap': True}
```

```
# fit a model with the parameters that were found so that I can get the feature impor
```

```
RF_test = RandomForestClassifier(n_estimators=1000, min_samples_split=10, min_samples
```

```
RF_test.fit(X_res_train, y_res_train)
```

```
y_pred=RF_test.predict(X_test)
```

```
print(classification_report(y_test, y_pred))
```

```
# so somehow with this random state the performance was just slightly better than be
```

	precision	recall	f1-score	support
0	0.79	0.84	0.81	2696
1	0.79	0.72	0.75	2232
micro avg	0.79	0.79	0.79	4928
macro avg	0.79	0.78	0.78	4928
weighted avg	0.79	0.79	0.79	4928

```
In [214]: # Get final model feature importance values and then plot
```

```
names = X.columns
```

```
feature_imp = pd.Series(RF_test.feature_importances_, index=names).sort_values(ascend
```

```
feature_imp
```

```
Out[214]: start_year      0.874300  
         salary          0.047493
```

```

start_month      0.032968
seniority         0.022776
dept_engineer    0.006700
dept_data_science 0.005563
dept_sales       0.004785
dept_marketing   0.003260
dept_design      0.002155
dtype: float64

```

```
In [215]: names = X.columns
```

```

feature_imp = pd.Series(RF_test.feature_importances_, index=names).sort_values(ascending=False)
feature_imp

```


```

plt.rcParams['figure.figsize'] = (15, 12)
plt.rcParams['axes.labelweight'] = 'bold'
#plt.rcParams["axes.labelsize"] = 30
rcParams['axes.titlepad'] = 20
sns.set(style='whitegrid')
sns.set_context("paper", rc={"font.size":32,"axes.titlesize":32,"axes.labelsize":30,

sns.barplot(x=feature_imp, y=feature_imp.index, palette='Reds_d')
# Add labels to your graph
plt.xlabel('Importance Score')
#plt.ylabel('Features')
plt.title("Visualizing Important Features from a Random Forest Classifier")

```

```
Out[215]: Text(0.5, 1.0, 'Visualizing Important Features from a Random Forest Classifier')
```



```
In [ ]: # OVERALL SUMMARY AND RESPONSES TO GOAL AND HINT QUESTIONS:
        # A random forest classifier was used to predict whether an employee is still working or not.

        #####

        #### Caveats:
        # I have tried includeing and excluding company ID in my predictive models. It did not work.
        # I also did not include number of days as a covariate or run a regression model with it.
        # In most of my analyses here, I used start year to potentiall explain stochasticity and variance.

        #####

        #### The goal was to predict employee retention and understand its main drivers.
        # According to the RF model the main drivers of whether an employee was retained or not were:
        # In terms of salary: Employees at medium and high salary levels (compared to low and high).
        # In terms of seniority: Again, employees within the mid-range seniority category (more than 10 years).
        # In terms of start month: Only real anomoly is November; people whose hire date falls in November.

        # If you look at the further exploratory analysis below, you can see that really, year of hire is a strong predictor.

        #####

        #### Further questions:
```

```

## What are the main factors that drive employee churn? Do they make sense?
# We really need more information, but it seems like incentives for mid-range salary a

## What might you be able to do for the company to address employee Churn, what would
# Ask for more information about employees themselves. We do not have data from very f

## If you could add to this data set just one variable that could help explain employe
# I would want some information about rate of change in salary as well as some informa

```

```
In [ ]: ##### FURTHER INVESTIGATION BELOW (EXTRA) ##### -->
```

```
In [229]: #dc1_NEW.head(20)
```

```
In [218]: # Further exploration:
# RF model with company included and all 12 companies.
# Here the outcome/predicted variable will be binary, still employeeed or not.
y = dc1_NEW[['still_employeeed']]

# Here I use the full company set (1-12) company ID.
X = dc1_NEW[['company_id','dept','seniority', 'salary', 'start_year', 'start_month']]

In [219]: # Get dummy variables for categorical features; use drop_first to get rid of the ext
X = pd.concat([X,pd.get_dummies(X['dept'], prefix='dept', drop_first=True)],axis=1)
X = pd.concat([X,pd.get_dummies(X['company_id'], prefix='company', drop_first=True)]

# drop old dept column since it is no longer informative.
X = X.drop(['dept'], axis=1)
X = X.drop(['company_id'], axis=1)
X.head()
```

```
Out[219]:
```

	seniority	salary	start_year	start_month	dept_data_science	\
0	28	89000.0	2014	3	0	
1	20	183000.0	2013	4	0	
2	14	101000.0	2014	10	0	
3	20	115000.0	2012	5	0	
4	23	276000.0	2011	10	1	

	dept_design	dept_engineer	dept_marketing	dept_sales	company_2	\
0	0	0	0	0	0	
1	0	0	1	0	0	
2	0	0	1	0	0	
3	0	0	0	0	0	
4	0	0	0	0	1	

	company_3	company_4	company_5	company_6	company_7	company_8	\
0	0	0	0	0	1	0	
1	0	0	0	0	1	0	
2	0	1	0	0	0	0	

3	0	0	0	0	1	0
4	0	0	0	0	0	0

	company_9	company_10	company_11	company_12
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

```
In [220]: # create training and test data sets based on X and y defined above. Will set test size to 0.2
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42)
```

```
In [221]: # check for class imbalance
y_train.value_counts()
```

```
Out[221]: 0    10801
          1     8941
          Name: still_employed, dtype: int64
```

```
In [222]: sm = SMOTE(random_state=11)
X_res_train, y_res_train = sm.fit_resample(X_train, y_train)
print('Resampled dataset shape %s' % Counter(y_res_train))
```

```
Resampled dataset shape Counter({'0': 10801, '1': 10801})
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/utils/validation.py:761: DataConversionWarning:
y = column_or_1d(y, warn=True)
```

```
In [225]: # just try an out-of-the box RF here to see how it performs compared to previous models
```

```
RF = RandomForestClassifier()
RF.fit(X_res_train, y_res_train)
y_pred=RF.predict(X_test)
print(classification_report(y_test, y_pred))
# performance is about equal - maybe a little bit better, but lots more features sin
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/ensemble/forest.py:246: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

	precision	recall	f1-score	support
0	0.76	0.84	0.80	2700
1	0.78	0.68	0.72	2236
micro avg	0.76	0.76	0.76	4936

macro avg	0.77	0.76	0.76	4936
weighted avg	0.77	0.76	0.76	4936

```
In [226]: print('Overall accuracy of a binary RF classifier on the Test Set:', metrics.accuracy_score(y_test, y_pred))
```

Overall accuracy of a binary RF classifier on the Test Set: 0.7647893030794165

```
In [227]: # Get feature importance values and then plot
names = X.columns
feature_imp = pd.Series(RF.feature_importances_, index=names).sort_values(ascending=False)
feature_imp
```

```
Out[227]: start_year      0.454974
salary      0.204059
seniority    0.129628
start_month  0.102820
company_2    0.013910
dept_marketing 0.009270
company_3    0.009185
dept_engineer 0.008783
company_4    0.007978
dept_sales   0.007902
company_5    0.007783
dept_data_science 0.007013
company_6    0.006659
company_7    0.006651
dept_design  0.005853
company_10   0.005708
company_9    0.005627
company_8    0.005559
company_12   0.000321
company_11   0.000315
dtype: float64
```

```
In [231]: names = X.columns


feature_imp = pd.Series(RF.feature_importances_, index=names).sort_values(ascending=False)
feature_imp

plt.rcParams['figure.figsize'] = (10, 10)
plt.rcParams['axes.labelweight'] = 'bold'
rcParams['axes.titlepad'] = 20
sns.set(style='whitegrid')
sns.set_context("paper", rc={"font.size":32,"axes.titlesize":32,"axes.labelsize":30,

sns.barplot(x=feature_imp, y=feature_imp.index, palette='Blues_d')
```

```
# Add labels to your graph
plt.xlabel('Importance Score')
#plt.ylabel('Features')
plt.title("Visualizing Important Features from a Random Forest Classifier")
```

Out[231]: Text(0.5, 1.0, 'Visualizing Important Features from a Random Forest Classifier')

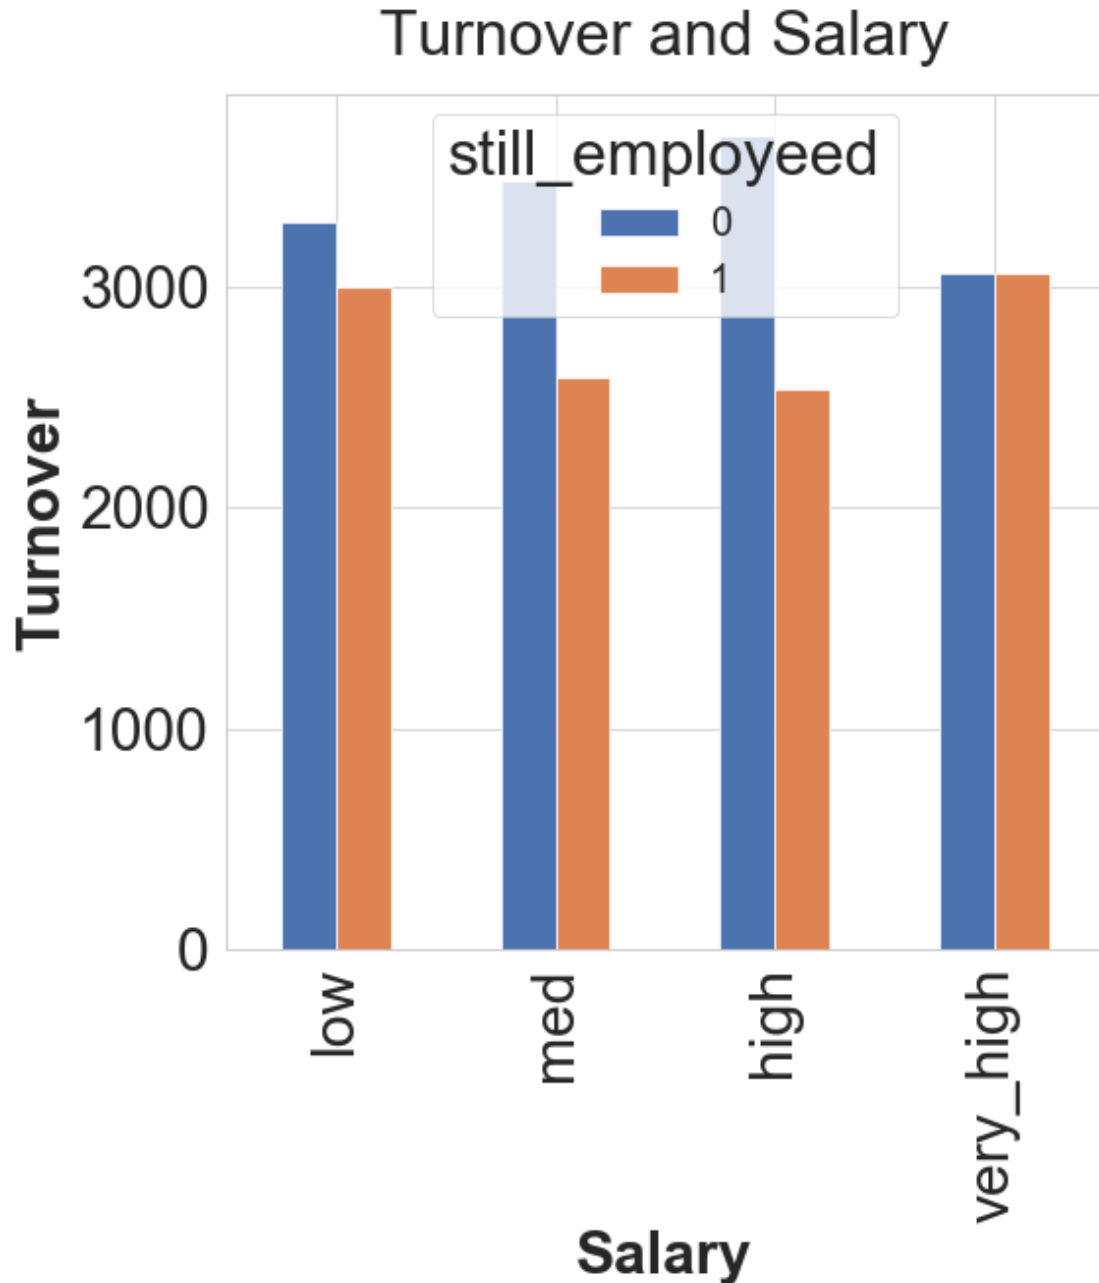


In []: *# Overall, same features are important -- start year appears to be really driving trend*
Salary and seniority appear to be the next most important. Let's look closer since t
If I had more time I could look at partial dependence plots to isolate effect of fea

In [241]: *# Reminder: we saw this fig before -- maybe there is something here (salary has been*
Perhaps at the highest salary level, it's true that equal numbers of people are sta
At medium and high salary levels, more people are leaving than going. Maybe they
The gap between those who quit and those who are still working is larger for peopl

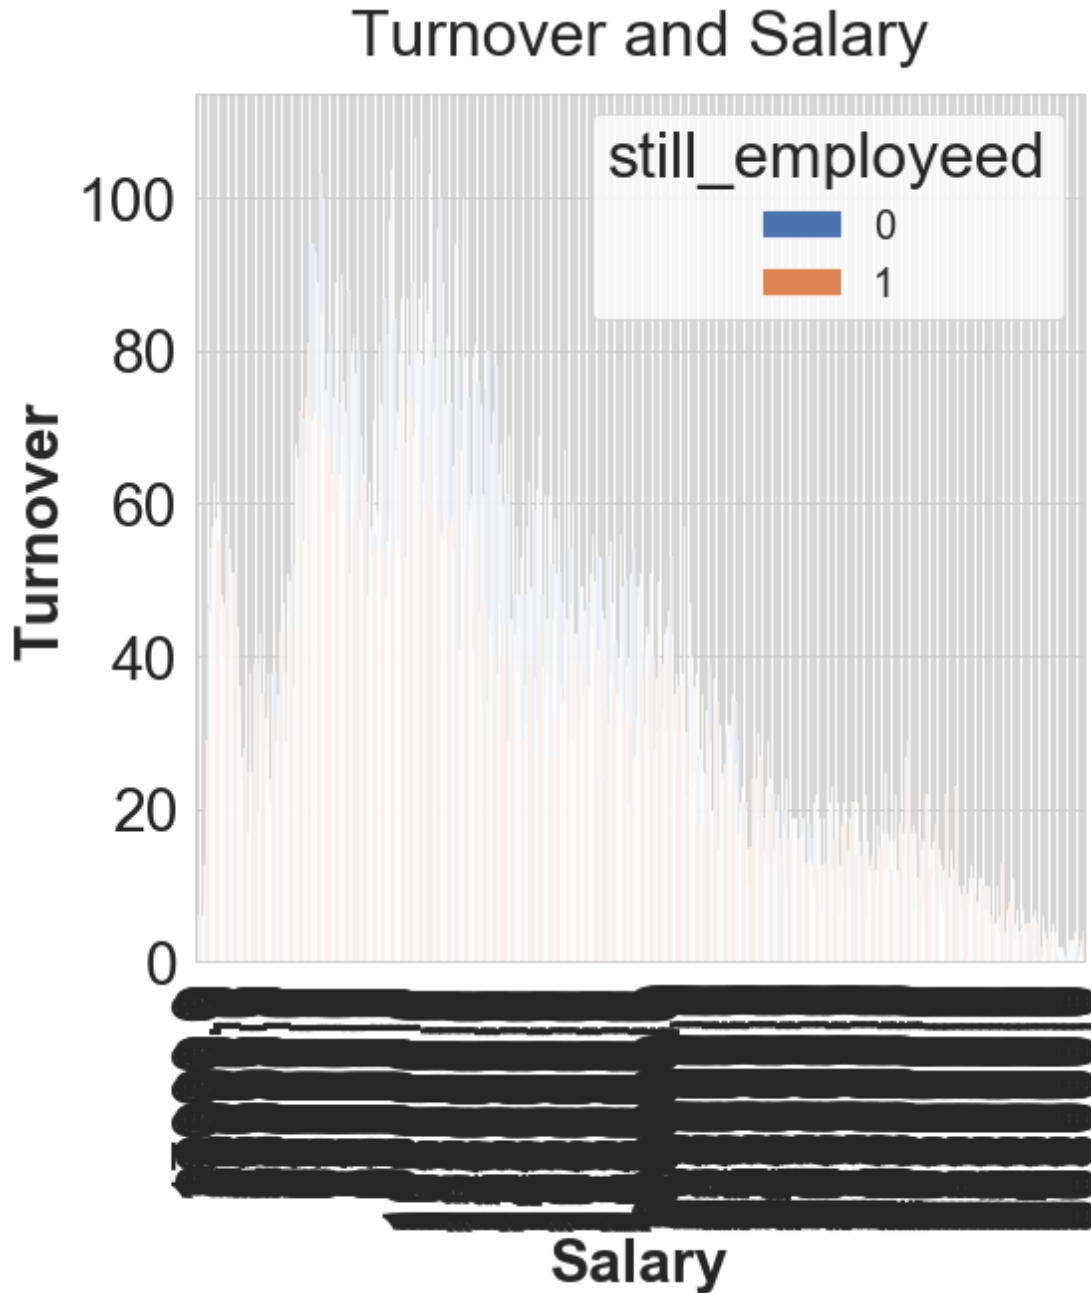
```
# Salary and turnover classification (still employed or not)
plt.rcParams['figure.figsize'] = (8, 8)
pd.crosstab(dc1_NEW.salary_rank, dc1_NEW.still_employed).plot(kind='bar')
plt.title('Turnover and Salary')
plt.xlabel('Salary')
plt.ylabel('Turnover')
```


Out[241]: Text(0, 0.5, 'Turnover')



```
In [267]: # Salary and turnover classification (still employed or not)
plt.rcParams['figure.figsize'] = (8, 8)
pd.crosstab(dc1_NEW.salary, dc1_NEW.still_employed).plot(kind='bar')
plt.title('Turnover and Salary')
plt.xlabel('Salary')
plt.ylabel('Turnover')
```

```
Out[267]: Text(0, 0.5, 'Turnover')
```



```
In [ ]: # Here are a couple of new visualizations for start year and month, with these effects
```

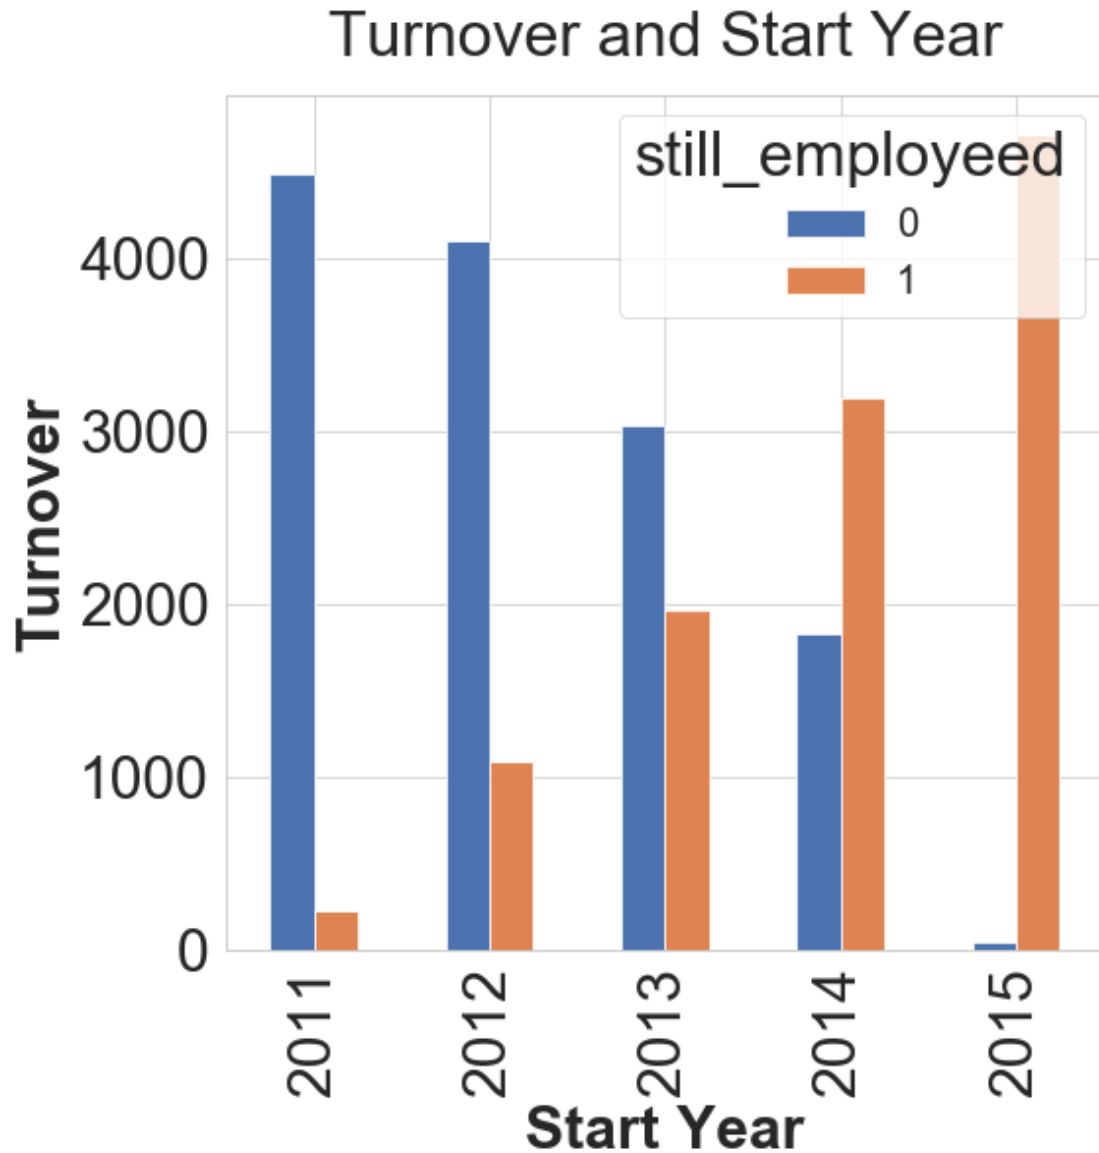
```
In [242]: pd.crosstab(dc1_NEW.start_year,dc1_NEW.still_employed).plot(kind='bar')
plt.title('Turnover and Start Year')
plt.xlabel('Start Year')
plt.ylabel('Turnover')
```

```
# CONCLUDE:
```

```
# Including start year is sort of like an artifact. More people are likely to have q
```

```
# Perhaps this should really be removed from the model, because it's giving the 'wro
```

```
Out[242]: Text(0, 0.5, 'Turnover')
```




```
In [243]: pd.crosstab(dc1_NEW.start_month,dc1_NEW.still_employed).plot(kind='bar')
plt.title('Turnover and Start Month')
plt.xlabel('Start Month')
plt.ylabel('Turnover')
```

```
# CONCLUDE
```

```
# The only month where number of people who have stayed vs. left is higher is Novemb
```

```
Out [243]: Text(0, 0.5, 'Turnover')
```



```
In [244]: # One last thing... what happens when start year is removed from the final model above?
```

```
y = dc1_NEW[['still_employed']]
```

```
# Here I use the full company set (1-12) company ID.
```

```
X = dc1_NEW[['company_id', 'dept', 'seniority', 'salary', 'start_month']]
```

```
In [245]: # Get dummy variables for categorical features; use drop_first to get rid of the extra
```

```
X = pd.concat([X, pd.get_dummies(X['dept'], prefix='dept', drop_first=True)], axis=1)
```

```
X = pd.concat([X, pd.get_dummies(X['company_id'], prefix='company', drop_first=True)], axis=1)
```

```
# drop old dept column since it is no longer informative.
```

```
X = X.drop(['dept'], axis=1)
X = X.drop(['company_id'], axis=1)
X.head()
```

```
Out [245]:
```

	seniority	salary	start_month	dept_data_science	dept_design	\
0	28	89000.0	3	0	0	
1	20	183000.0	4	0	0	
2	14	101000.0	10	0	0	
3	20	115000.0	5	0	0	
4	23	276000.0	10	1	0	

	dept_engineer	dept_marketing	dept_sales	company_2	company_3	company_4	\
0	0	0	0	0	0	0	
1	0	1	0	0	0	0	
2	0	1	0	0	0	1	
3	0	0	0	0	0	0	
4	0	0	0	1	0	0	

	company_5	company_6	company_7	company_8	company_9	company_10	\
0	0	0	1	0	0	0	
1	0	0	1	0	0	0	
2	0	0	0	0	0	0	
3	0	0	1	0	0	0	
4	0	0	0	0	0	0	

	company_11	company_12
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

```
In [247]: # create training and test data sets based on X and y defined above. Will set test s
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_st
```

```
In [248]: # check for class imbalance
y_train.still_employeed.value_counts()
```

```
Out [248]:
```

0	10801
1	8941

Name: still_employeed, dtype: int64

```
In [249]: sm = SMOTE(random_state=11)
X_res_train, y_res_train = sm.fit_resample(X_train, y_train)
print('Resampled dataset shape %s' % Counter(y_res_train))
```

```
Resampled dataset shape Counter({'0': 10801, '1': 10801})
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/utils/validation.py:761: DataConversionWarning:
  y = column_or_1d(y, warn=True)
```

```
In [250]: # just try an out-of-the box RF here to see how it performs compared to previous mod
```

```
RF = RandomForestClassifier()
RF.fit(X_res_train, y_res_train)
y_pred=RF.predict(X_test)
print(classification_report(y_test, y_pred))
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/ensemble/forest.py:246: FutureWarning: The default
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

	precision	recall	f1-score	support
0	0.56	0.64	0.60	2700
1	0.47	0.38	0.42	2236
micro avg	0.53	0.53	0.53	4936
macro avg	0.51	0.51	0.51	4936
weighted avg	0.52	0.53	0.52	4936

```
In [251]: print('Overall accuracy of a binary RF classifier on the Test Set:', metrics.accuracy
```

```
Overall accuracy of a binary RF classifier on the Test Set: 0.5251215559157212
```

```
In [252]: # Get feature imporance values and then plot
```

```
names = X.columns
feature_imp = pd.Series(RF.feature_importances_, index=names).sort_values(ascending=F
feature_imp
```

```
Out[252]: salary          0.435520
seniority          0.250622
start_month        0.162833
company_2          0.020925
company_3          0.013500
dept_engineer      0.012430
company_4          0.011655
dept_sales         0.011543
company_5          0.010378
dept_marketing     0.009656
company_8          0.009457
company_7          0.009100
dept_data_science 0.009088
```

```

company_6          0.008901
dept_design        0.008119
company_10         0.007977
company_9          0.007313
company_11         0.000523
company_12         0.000462
dtype: float64

```

```
In [253]: names = X.columns
```

```

feature_imp = pd.Series(RF.feature_importances_,index=names).sort_values(ascending=False)
feature_imp

```


```

plt.rcParams['figure.figsize'] = (10, 10)
plt.rcParams['axes.labelweight'] = 'bold'
rcParams['axes.titlepad'] = 20
sns.set(style='whitegrid')
sns.set_context("paper", rc={"font.size":32,"axes.titlesize":32,"axes.labelsize":30,

sns.barplot(x=feature_imp, y=feature_imp.index, palette='Greens_d')
# Add labels to your graph
plt.xlabel('Importance Score')
#plt.ylabel('Features')
plt.title("Visualizing Important Features from a Random Forest Classifier")

```

```
Out[253]: Text(0.5, 1.0, 'Visualizing Important Features from a Random Forest Classifier')
```



```
In [283]: #Just for fun
sns.boxplot(x="still_employeed", y="salary", hue = 'seniority', data=dc1_NEW)
# Put the legend out of the figure
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)

Out[283]: <matplotlib.legend.Legend at 0x1a2c662da0>
```

