

Vladimir_Shteyn_Yammer

February 27, 2019

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import os
import pandas as pd

from insight_datachallenges.week2 import DATADIR
```

0.1 Overall approach

I was provided with a chart showing user engagement over calendar time, as well as server log data on each user. I noticed right away that user engagement dropped at the end of July. Is this drop significant? And, what could be the reason for it?

0.2 Conclusion

Vacation time in Europe is the likely cause of the decrease in user engagement.

```
In [2]: events_csv_filename = os.path.join(DATADIR, 'yammer_events.csv')
events = pd.read_csv(events_csv_filename)
```

I wasn't sure whether user engagement events were logged by user's local time or a time of day at the server's location. To infer this, I plotted the number of users signing up for Yammer as a function of time of day. If events were logged user-local time, I'd expect them to be distributed during typical waking hours. (This assumes a uniform distribution of users across the world).

```
In [3]: events['occurred_at'] = pd.to_datetime(events['occurred_at'])
events['occurred_at_date'] = events['occurred_at'].apply(lambda x: x.date())
events['time_of_day'] = events['occurred_at'].apply(lambda x: x.hour)
```

```
In [4]: engagement_mask = events['event_type'] == 'signup_flow'
```

```
In [5]: event_time_groupby = events.loc[engagement_mask, :].groupby(['time_of_day'])
```

```
In [6]: # a peek at the data (sanity check)
events.head()
```

```
Out[6]:
```

	user_id	occurred_at	event_type	event_name	location	\
0	10522.0	2014-05-02 11:02:39	engagement	login	Japan	


1	10522.0	2014-05-02	11:02:53	engagement	home_page	Japan
2	10522.0	2014-05-02	11:03:28	engagement	like_message	Japan
3	10522.0	2014-05-02	11:04:09	engagement	view_inbox	Japan
4	10522.0	2014-05-02	11:03:16	engagement	search_run	Japan

			device	user_type	occurred_at_date	time_of_day
0	dell	inspiron	notebook	3.0	2014-05-02	11
1	dell	inspiron	notebook	3.0	2014-05-02	11
2	dell	inspiron	notebook	3.0	2014-05-02	11
3	dell	inspiron	notebook	3.0	2014-05-02	11
4	dell	inspiron	notebook	3.0	2014-05-02	11

From the plot below, it's clear sign ups occur during normal waking hours, so the occurred_at column is indeed user local time.

```
In [7]: counts_time_of_day = event_time_groupby.agg('count')
# df = counts_time_of_day.unstack(level=0)
fig, ax = plt.subplots(1, 1, figsize=(8, 8))
ax.set_ylabel('count', fontsize=18)
ax.set_xlabel('time of day (24 hour scale)', fontsize=18)
counts_time_of_day.loc[:, 'user_id'].plot(ax=ax)
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x7f87aed75470>
```



Let's just confirm the assumption that users are evenly distributed around the world. Below I've printed out the number of user events for each country.

```
In [8]: location_groupby = events.groupby('location')
        location_groupby.agg('count')['user_id']
```

```
Out[8]: location
Argentina          1717
Australia          7494
Austria            2494
Belgium            2822
Brazil             11240
Canada              9126
Chile              1092
Colombia           1945
Denmark            2191
```

Egypt	2258
Finland	1926
France	17364
Germany	23524
Greece	989
Hong Kong	1525
India	9620
Indonesia	6224
Iran	3122
Iraq	1400
Ireland	1073
Israel	2130
Italy	11790
Japan	26046
Korea	7180
Malaysia	2529
Mexico	9106
Netherlands	4494
Nigeria	1642
Norway	2020
Pakistan	1035
Philippines	1373
Poland	3803
Portugal	1336
Russia	12226
Saudi Arabia	4104
Singapore	1497
South Africa	2324
Spain	5874
Sweden	3901
Switzerland	3760
Taiwan	3600
Thailand	2008
Turkey	2432
United Arab Emirates	2343
United Kingdom	16475
United States	94728
Venezuela	1930

Name: user_id, dtype: int64

It looks like most users are from Europe, which suggests the end-of-July drop in engagement is due to Europeans going on holiday the month of August.

In []: